

# **Computational Analysis of Educational Data: A Field Guide Using R**

Wei Wang, Mete Akcaoglu, Joshua Rosenberg, Shaun Kellogg

2025-10-16

# Table of contents

<b>Preface</b>	<b>5</b>
<b>Preface</b>	<b>6</b>
<b>Section 1 Getting Started</b>	<b>7</b>
<b>Section 2 Capturing and Analyzing Text Data with Computational Methods</b>	<b>8</b>
2.1 Overview . . . . .	8
2.2 Accessing Text Data . . . . .	8
2.2.1 Web Scraping (Unstructured or API) . . . . .	9
2.2.2 Audio Transcripts (Zoom, etc.) . . . . .	12
2.2.3 PDF . . . . .	13
2.2.4 Survey, Discussions, etc. . . . .	16
2.3 Frequency-based Analysis . . . . .	18
2.3.1 Purpose . . . . .	18
2.3.2 Sample Research Questions . . . . .	18
2.3.3 Sample Methods . . . . .	19
2.3.4 Analysis . . . . .	19
2.3.5 Results and Discussions . . . . .	22
2.4 Dictionary-based Analysis . . . . .	23
2.4.1 Purpose . . . . .	23
2.4.2 Sample Research Questions . . . . .	23
2.4.3 Analysis . . . . .	23
2.4.4 Results and Discussions . . . . .	26
2.5 Clustering-Based Analysis . . . . .	26
2.5.1 Purpose . . . . .	27
2.5.2 Sample Research Questions . . . . .	27
2.5.3 Analysis . . . . .	27
2.5.4 Results and Discussions . . . . .	31
<b>Section 3 Social Network Analyses (Relational Data)</b>	<b>34</b>
3.1 Overview . . . . .	34
3.2 Accessing SNA Data . . . . .	34
3.2.1 Types and Sources of SNA Data . . . . .	34
3.2.2 Example 1: Creating a Simple Network from an Edge List . . . . .	35
3.2.3 Example 2: Generating Network Data from Digital Traces . . . . .	35

3.2.4 Example 3: Collecting SNA Data via Surveys . . . . .	35
3.2.5 Node Attribute Data . . . . .	36
3.2.6 Further Examples . . . . .	36
3.2.7 Best Practices and Tips . . . . .	37
3.2.8 Summary . . . . .	37
3.3 Network Management & Measurement in Social Network Analysis . . . . .	37
3.3.1 Purpose + Case . . . . .	37
3.3.2 Sample Research Questions . . . . .	38
3.3.3 Analysis . . . . .	38
3.3.4 Results and Discussion . . . . .	48
3.4 Case Study: Hashtag Common Core . . . . .	50
3.4.1 Purpose & Case . . . . .	50
3.4.2 Sample Research Questions . . . . .	50
3.4.3 Analysis . . . . .	51
3.4.4 Results and Discussion . . . . .	57
<b>Section 4 Secondary Analysis of Big Data (Numeric Data)</b>	<b>61</b>
4.1 Overview . . . . .	61
4.2 Accessing Big data (Broadening the Horizon) . . . . .	61
4.2.1 Big Data . . . . .	61
4.2.2 Learning Analytics . . . . .	67
4.3 Logistic Regression ML . . . . .	70
4.3.1 Purpose + CASE . . . . .	70
4.3.2 Sample Research Questions (RQs) . . . . .	71
4.3.3 Analysis . . . . .	71
4.3.4 Results and Discussions . . . . .	75
4.4 Random Forests ML on Interactions Data . . . . .	77
4.4.1 Purpose + CASE . . . . .	77
4.4.2 Sample Research Questions . . . . .	77
4.4.3 Analysis . . . . .	78
4.4.4 Results and Discussions . . . . .	85
<b>Section 5 Multimodal Data (Images, Video, Audio) with LLMs</b>	<b>87</b>
5.1 Overview . . . . .	87
5.2 Images . . . . .	87
5.3 Analyzing Images . . . . .	88
5.3.1. Setting Up Kuzco . . . . .	88
5.3.2 Image Classification . . . . .	89
Case Study: Analyzing Classroom Photographs with Kuzco to Explore Student Engagement . . . . .	89
5.3.2.2 Research Questions . . . . .	89
5.3.2.3 Methods . . . . .	90
5.3.2.4 Data Source . . . . .	90

5.3.2.5 Data Analysis . . . . .	90
5.3.2.6 Results and Discussion: . . . . .	94
<b>Section 6 Local LLMs in Educational Research</b>	<b>96</b>
6.1 What are Local LLMs? . . . . .	96
6.2 What Can Local LLMs Do? . . . . .	96
6.3 Getting Started with LM Studio . . . . .	97
6.3.1 Installation Steps . . . . .	97
6.3.2 Main Features . . . . .	97
6.3.3 API Integration . . . . .	98
6.3.4 Summary Table of LM Studio Capabilities: . . . . .	98
6.4 Case Study: Comparing Local LLM Analysis to Traditional NLP on University AI Policy Texts . . . . .	99
6.4.1 Research Question . . . . .	99
6.4.2 Data Context . . . . .	99
6.4.3 Implementation with LM Studio (Thematic Analysis) . . . . .	100
6.4.4 Sample Output . . . . .	106
6.4.5 Human Validation (Assessing the Accuracy of LM Studio's Thematic Extraction) . . . . .	111
6.4.5.3 Case Study Discussion . . . . .	121
6.4.6 Reflection . . . . .	123
<b>Section 7 Communication Collaboration Practices</b>	<b>125</b>
Abstract . . . . .	125
<b>Conclusion</b>	<b>126</b>
<b>References</b>	<b>127</b>

# Preface

# Preface

Welcome to *Computational Analysis of Educational Data: A Field Guide Using R!*

The book is organized around six sections. Within these seven sections are specific chapters, which serve as cookbook “entries”. While the section overviews (the first bullet point within each section) introduce the techniques or methodologies used in the section’s chapters, the entries (the subsequent bullet points) are intended to address a specific, narrow problem, as well as to provide a sample for researchers in writing their research questions, methods, results (and discussions) sections based on the analyses.

# Section 1 Getting Started

**Abstract:** In this section, we will introduce the steps to get started with using R and RStudio. The section will also cover the key concepts in working with R such as packages, R markdown. We will also cover important steps in preparing data, which we will be using frequently in the rest of the book. Since these topics have been covered in so much more detail in other books, or other resources online, we will end the section with a list of key resources on these topics that are easily accessible online.

- Section Overview
- Using R and RStudio
  - Foundational skills: Data, packages, functions, projects
  - Reproducibility and R Markdown documents
  - Base R: assignment, equality, dollar-sign indexing, xx
  - R files
  - Project setup within RStudio
- Core Data Wrangling Capabilities
  - Selecting and renaming variables
  - Creating new variables
  - Arranging variables
  - Joining data

# **Section 2 Capturing and Analyzing Text Data with Computational Methods**

## **2.1 Overview**

In social sciences, analyzing text data is usually considered the “job” of qualitative researchers. Traditionally, qualitative research involves identifying patterns in non-numeric data, and this pattern recognition is typically done manually. This process is time-intensive but can yield rich research results. Traditional methods for analyzing text data involve human coding and can include direct sources (e.g., books, online texts) or indirect sources (e.g., interview transcripts).

With the advent of new software, we can capture and analyze text data in ways that were previously not possible. Modern data collection techniques include web scraping, accessing social media APIs, or downloading large online documents. Given the increased size of text data, analysis now requires computational approaches (e.g., dictionary-based, frequency-based, machine learning) that go beyond manual coding. These computational methods allow social scientists to ask new types of research questions, expanding the scope and depth of possible insights.

**Disclaimer:** While resources are available that discuss these analysis methods in depth, this book aims to provide a practical guide for social scientists, using data they will likely encounter. Our goal is to present a “cookbook” for guiding research projects through real-world examples.

## **2.2 Accessing Text Data**

For conventional qualitative researcher, text data comes from the usual sources such as interview transcripts or existing documents. Nowadays, however, text can be found and collected in many different ways. For example, social media is rich with text data, likewise for faculty who are teaching course (especially online), every student writing can be seen as a piece of text data. In this section, we will cover a few basic ways of accessing text data. Please note that although we cover some prominent ways, this is by no means an exhaustive list. Therefore, please refer to the additional resources section at the end of the section to dive deeper.

## 2.2.1 Web Scraping (Unstructured or API)

### What is Web Scraping?

Web scraping refers to the automated process of extracting data from web pages. It is particularly useful when dealing with extensive lists of websites that would be tedious to mine manually. A typical web scraping program follows these steps:

1. Loads a webpage.
2. Downloads the HTML or XML structure.
3. Identifies the desired data.
4. Converts the data into a format suitable for analysis, such as a data frame.

In addition to text, web scraping can also be used to download other content types, such as audio-visual files.

### Is Web Scraping Legal?

Web scraping was common in the early days of the internet, but with the increasing value of data, legal norms have evolved. To avoid legal issues, check the “Terms of Service” for specific permissions on the website, often accessible via “robots.txt” files. Consult legal advice when in doubt.

### Reading a Web Page into R

Once permissions are confirmed, the first step in web scraping is to download the webpage’s source code into R, typically using the `rvest` package by Hadley Wickham.

```
# Install and load the rvest package
install.packages("rvest")
library(rvest)
```

To demonstrate, we will scrape a simple Wikipedia page. Static pages, which lack interactive elements like JavaScript, are simpler to scrape. You can view the page’s HTML source in your browser by selecting **Developer Tools > View Source**.

```
# Load the webpage
wikipedia_page <- read_html("https://en.wikipedia.org/wiki/World_Health_Organization_ranking")

# Verify that the webpage loaded successfully
wikipedia_page
```

## Parsing HTML

The next challenge is extracting specific information from the HTML structure. HTML files have a “tree-like” format, allowing us to target particular sections. Use your browser’s “Developer Tools” to inspect elements and locate the data. Right-click the desired element and select **Inspect** to view its structure.

To isolate data sections within the HTML structure, identify the **XPath** or **CSS selectors**. For instance:

```
# Extract specific section using XPath
section_of_wikipedia <- html_node(wikipedia_page, xpath='//*[@id="mw-content-text"]/div/table')
head(section_of_wikipedia)
```

To convert the extracted section into a data frame, use **html\_table()**:

```
# Convert the extracted data into a table
health_rankings <- html_table(section_of_wikipedia)
head(health_rankings[, (1:2)]) # Display the first two columns
```

## Parsing with CSS Selectors

For more complex web pages, CSS selectors can be an alternative to XPath. Tools like **Selector Gadget** can help identify the required CSS selectors.

For example, to scrape event information from Duke University’s main page:

```
# Load the webpage
duke_page <- read_html("https://www.duke.edu")

# Extract event information using CSS selector
duke_events <- html_nodes(duke_page, css="li:nth-child(1) .epsilon")
html_text(duke_events)
```

## Scraping with Selenium

For tasks involving interactive actions (e.g., filling search fields), use **RSelenium**, which enables automated browser operations.

To set up Selenium, install the **Java SE Development Kit** and **Docker**. Then, start Selenium in R:

```

# Install and load RSelenium
install.packages("RSelenium")
library(RSelenium)

# Start a Selenium session
rD <- rsDriver()
remDr <- rD$client
remDr$navigate("https://www.duke.edu")

```

To automate data entry, identify the CSS selector for the search box and input the query:

```

# Find the search box element and enter a query
search_box <- remDr$findElement(using = 'css selector', 'fieldset input')
search_box$sendKeysToElement(list("data science", "\uE007")) # "\uE007" represents Enter key

```

## **Web Scraping within a Loop**

To scrape multiple pages, embed the code within a loop to automate tasks across different URLs. Since each site may have a unique structure, generalized scraping can be time-intensive and error-prone. Implement error handling to manage interruptions.

## **When to Use Web Scraping**

Web scraping is appropriate if:

- **Page structure is consistent across sites:** For example, a government site with date suffixes but a uniform layout.
- **Manual data collection is prohibitive:** For extensive text or embedded tables.

When feasible, consider alternatives like APIs or data-entry services (e.g., Amazon Mechanical Turk) for better efficiency and legal compliance.

## **What is an API?**

An Application Programming Interface (API) is a set of protocols that allows computers to communicate and exchange information. A common type is the REST API, where one machine sends a request, and another returns a response. APIs provide standardized access to data, services, and functionalities, making them essential in software development.

## When to Use an API

APIs are commonly used for:

- **Integrating with Third-Party Services:** APIs connect applications to services like payment gateways or social media.
- **Accessing Data:** APIs retrieve data from systems or databases (e.g., real-time weather data).
- **Automating Tasks:** APIs automate processes within applications, such as email marketing.
- **Building New Applications:** APIs allow developers to build new apps or services (e.g., a mapping API for navigation).
- **Streamlining Workflows:** APIs enable seamless communication and data exchange across systems.

## Using Reddit API with RedditExtractoR

Reddit is a social media platform featuring a complex network of users and discussions, organized into “subreddits” by topic. RedditExtractoR, an R package, enables data extraction from Reddit to identify trends and analyze interactions.

```
# Install and load RedditExtractoR
install.packages("RedditExtractoR")
library(RedditExtractoR)

# Access data from the GenAI subreddit
GenAI_reddit <- find_thread_urls(subreddit = "GenAI", sort_by = "new", period = "day")
view(GenAI_reddit)
```

### 2.2.2 Audio Transcripts (Zoom, etc.)

Audio transcripts are a rich source of text data, especially useful for capturing spoken content from meetings, interviews, or webinars. Many platforms, such as Zoom, provide automated transcription services that can be downloaded as text files for analysis. By processing these transcripts, researchers can analyze conversation themes, sentiment, or other linguistic features. Here’s how to access and prepare Zoom transcripts for analysis in R.

## Key Steps for Extracting Text Data from Audio Transcripts

### 1. Access the Zoom Transcript

- Log in to your Zoom account.
- Navigate to the “Recordings” section.
- Select the recording you wish to analyze and download the “Audio Transcript” file.

### 2. Import the Transcript into R

Once the file is downloaded, you can load it into R for analysis. Depending on the file format (usually a .txt file with tab or comma delimiters), use `read.table()`, `read.csv()`, or functions from the `readr` package to load the data.

```
# Load the transcript into R
transcript_data <- read.table("path/to/your/zoom_transcript.txt", sep = "\t", header = TRUE)
```

Adjust the `sep` parameter based on the delimiter used in the transcript file (typically \t for tab-delimited files).

### 3. Data Cleaning (if necessary)

Clean up the text data to remove unnecessary characters, standardize formatting, and prepare it for further analysis.

#### • Remove Unwanted Characters

Use `gsub()` to eliminate special characters and punctuation, keeping only alphanumeric characters and spaces.

```
# Remove special characters
transcript_data$text <- gsub("[^a-zA-Z0-9 ]", "", transcript_data$text)
```

#### • Convert Text to Lowercase

Standardize text to lowercase for consistency in text analysis.

```
# Convert text to lowercase
transcript_data$text <- tolower(transcript_data$text)
```

### 2.2.3 PDF

PDF files are a valuable source of text data, often found in research publications, government documents, and industry reports. We’ll explore two main methods for extracting text from PDFs:

1. **Extracting from Local PDF Files:** This method involves accessing and parsing text from PDF files stored locally, providing tools and techniques to efficiently retrieve text data from offline documents.
2. **Downloading and Extracting PDF Files:** This approach covers downloading PDFs from online sources and extracting their text. This method is useful for scraping publicly available documents from websites or databases for research purposes.

### 3. PDF Data Extractor (PDE)

For more advanced PDF text extraction and processing, you can use the [PDF Data Extractor \(PDE\) package](#). This package provides tools for extracting text data from complex PDF documents, supporting additional customization options for text extraction. PDE is a R package that easily extracts information and tables from PDF files. The `PDE_analyzer_i()` performs the sentence and table extraction while the included `PDE_reader_i()` allows the user-friendly visualization and quick-processing of the obtained results.

## Steps for Extracting Text from Local PDF Files

### 1. Install and Load the `pdftools` Package

Start by installing and loading the `pdftools` package, which is specifically designed for reading and extracting text from PDF files in R.

```
install.packages("pdftools")
library(pdftools)
```

### 2. Read the PDF as a Text File

Use the `pdf_text()` function to read the PDF file into R as a text object. This function returns each page as a separate string in a character vector.

```
txt <- pdf_text("path/to/your/file.pdf")
```

### 3. Extract Text from a Specific Page

To access a particular page from the PDF, specify the page number in the text vector. For example, to extract text from page 24:

```
page_text <- txt[24] # page 24
```

### 4. Extract Rows into a List

If the page contains a table or structured text, use the `scan()` function to read each row as a separate element in a list. The `textConnection()` function converts the page text for processing.

```
rows <- scan(textConnection(page_text), what = "character", sep = "\n")
```

#### 5. Split Rows into Cells

To further parse each row, split it into cells by specifying the delimiter, such as whitespace (using "\s+"). This converts each row into a list of individual cells.

```
row <- unlist(strsplit(rows[24], "\s+")) # Example with the 24th row
```

### Steps for Downloading and Extracting Text from PDF Files

#### 1. Download the PDF from the Web

Use the `download.file()` function to download the PDF file from a specified URL. Set the mode to "wb" (write binary) to ensure the file is saved correctly.

```
link <- paste0(  
  "http://www.singstat.gov.sg/docs/",  
  "default-source/default-document-library/",  
  "publications/publications_and_papers/",  
  "cop2010/census_2010_release3/",  
  "cop2010sr3.pdf"  
)  
download.file(link, "census2010_3.pdf", mode = "wb")
```

#### 2. Read the PDF as a Text File

After downloading, read the PDF into R as a text object using the `pdf_text()` function from the `pdftools` package. Each page of the PDF will be stored as a string in a character vector.

```
txt <- pdf_text("census2010_3.pdf")
```

#### 3. Extract Text from a Specific Page

Access the desired page (e.g., page 24) by specifying the page number in the character vector.

```
page_text <- txt[24] # Page 24
```

#### 4. Extract Rows into a List

Use the `scan()` function to split the page text into rows, with each row representing a line of text in the PDF. This creates a list where each line from the page is an element.

```
rows <- scan(textConnection(page_text), what = "character", sep = "\n")
```

## 5. Loop Through Rows and Extract Data

Starting from a specific row (e.g., row 7), loop over each row. For each row:

- Split the text by spaces ("\\s+") using `strsplit()`.
- Convert the result to a vector with `unlist()`.
- If the third cell in the row is not empty, store the second cell as `name` and the third cell as `total`, converting it to a numeric format after removing commas.

```
name <- c()
total <- c()

for (i in 7:length(rows)) {
  row <- unlist(strsplit(rows[i], "\\s+"))
  if (!is.na(row[3])) {
    name <- c(name, row[2])
    total <- c(total, as.numeric(gsub(", ", "", row[3])))
  }
}
```

### 2.2.4 Survey, Discussions, etc.

Surveys and discussion posts are valuable sources of text data in social science research, providing insights into participants' perspectives, opinions, and experiences. These data sources often come from open-ended survey responses, online discussion boards, or educational platforms. Extracting and preparing text data from these sources can reveal recurring themes, sentiment, and other patterns that support both quantitative and qualitative analysis. Below are key steps and code examples for processing text data from surveys and discussions in R.

### Key Steps for Processing Survey and Discussion Text Data

#### 1. Load the Data

Survey and discussion data are typically stored in spreadsheet formats like CSV. Begin by loading this data into R for processing. Here, `readr` is used for reading CSV files with `read_csv()`.

```
# Install and load necessary packages
install.packages("readr")
library(readr)

# Load data
survey_data <- read_csv("path/to/your/survey_data.csv")
```

## 2. Extract Text Columns

Identify and isolate the relevant text columns. For example, if the text data is in a column named “Response,” you can create a new vector for analysis.

```
# Extract text data from the specified column
text_data <- survey_data$Response
```

## 3. Data Cleaning

Prepare the text data by cleaning it, removing any unnecessary characters, and standardizing the text. This includes removing punctuation, converting text to lowercase, and handling extra whitespace.

- **Remove Unwanted Characters**

Use `gsub()` from base R to remove any non-alphanumeric characters, retaining only words and spaces.

```
# Remove special characters
text_data <- gsub("[^a-zA-Z0-9 ]", "", text_data)
```

- **Convert to Lowercase**

Standardize the text by converting all characters to lowercase.

```
# Convert text to lowercase
text_data <- tolower(text_data)
```

- **Remove Extra Whitespace**

Remove any extra whitespace that may be left after cleaning.

```
# Remove extra spaces
text_data <- gsub("\\s+", " ", text_data)
```

## 4. Tokenization and Word Counting (Optional)

If further analysis is needed, such as frequency-based analysis, split the text into individual words (tokenization) or count the occurrence of specific words. Here, `dplyr` is used to organize the word counts.

```
# Install and load necessary packages
install.packages("dplyr")
library(dplyr)

# Tokenize and count words
word_count <- strsplit(text_data, " ") %>%
  unlist() %>%
  table() %>%
  as.data.frame()
```

## 2.3 Frequency-based Analysis

In the following section, we will provide a “recipe” for the social scientist interested in these new methods of analyzing text data to get you from the initial stages of getting the data to running the analyses and the write up. Often left out is also a research question that suits or requires a method. Since we have a data and method-centric approach here, we will backtrack and also provide a research question, so that you can model after it in your own work. Finally, we will provide a sample results and discussions section.

### 2.3.1 Purpose

The purpose of the frequency-based approach is to count the number of words as they appear in a text file, whether it is a collection of tweets, documents, or interview transcripts. This approach aligns with the frequency-coding method (e.g., Saldana) and can supplement human coding by revealing the most/least commonly occurring words, which can then be compared across dependent variables.

#### Case Study: Frequency-Based Analysis of GenAI USage Guidelines in Higher Education

As AI writing tools like ChatGPT become more prevalent, educators are working to understand how best to integrate them within academic settings, while many students and instructors remain uncertain about acceptable use cases. Our research into AI usage guidelines from the top 100 universities in North America aims to identify prominent themes and concerns in institutional policies regarding GenAI.

### 2.3.2 Sample Research Questions

To investigate the nature of AI use policies within higher education institutions, in this study, our research questions are:

- **RQ1:** What are the most frequently mentioned words in university GenAI writing usage policies?
- **RQ2:** Which keywords reflect common concerns or focal points related to GenAI writing usage in academic settings?

### 2.3.3 Sample Methods

#### Data Source

The dataset consists of publicly available AI policy texts from 100 universities(USA), the data has been downloaded and saved as a CSV file for analysis. —we might have to write more here to model how a research should be describing the data from its acquisition to use for research.

#### Data Analysis

In order to analyze the data we used xyz, —let's provide a sample write up for the researcher to adapt.

### 2.3.4 Analysis

#### Step 1: Load Required Libraries

Install and load libraries for data processing, visualization, and word cloud generation.

```
# Install necessary packages if not already installed
#install.packages(c("tibble", "dplyr", "tidytext", "ggplot2", "viridis","tm", "wordcloud" "wordcloud2", "webshot"))

# Load libraries
library(readr)
library(tibble)
library(dplyr)
library(tidytext)
library(ggplot2)
library(viridis)
library(tm)
library(wordcloud)
library(wordcloud2)
library(webshot)
```

## Step 2: Load Data

Read the CSV file containing policy texts from the top 100 universities.

```
# Load the dataset  
university_policies <- read_csv("University_GenAI_Policy_Stance.csv")
```

## Step 3: Tokenize Text and Count Word Frequency

Process the text data by tokenizing words, removing stop words, and counting word occurrences.

```
# Tokenize text, remove stop words, and count word frequencies  
word_frequency <- university_policies %>%  
  unnest_tokens(word, Stance) %>% # Tokenize the 'Policy_Text' column  
  anti_join(stop_words) %>%          # Remove common stop words  
  count(word, sort = TRUE)            # Count and sort words by frequency  
word_frequency
```

```
# A tibble: 1,108 x 2  
  word           n  
  <chr>      <int>  
1 ai             124  
2 students       120  
3 chatgpt        80  
4 tools           73  
5 academic       68  
6 policy          35  
7 assignments     34  
8 instructors      34  
9 integrity        33  
10 student         32  
# i 1,098 more rows
```

## Step 4: Create a Word Cloud

Generate a word cloud to visualize the frequency of words in a circular shape.

```
# Create and display the GenAI usage Stance wordcloud  
  
wordcloud(words = word_frequency$word, freq = word_frequency$n, scale = c(4, 0.5), random.or...
```



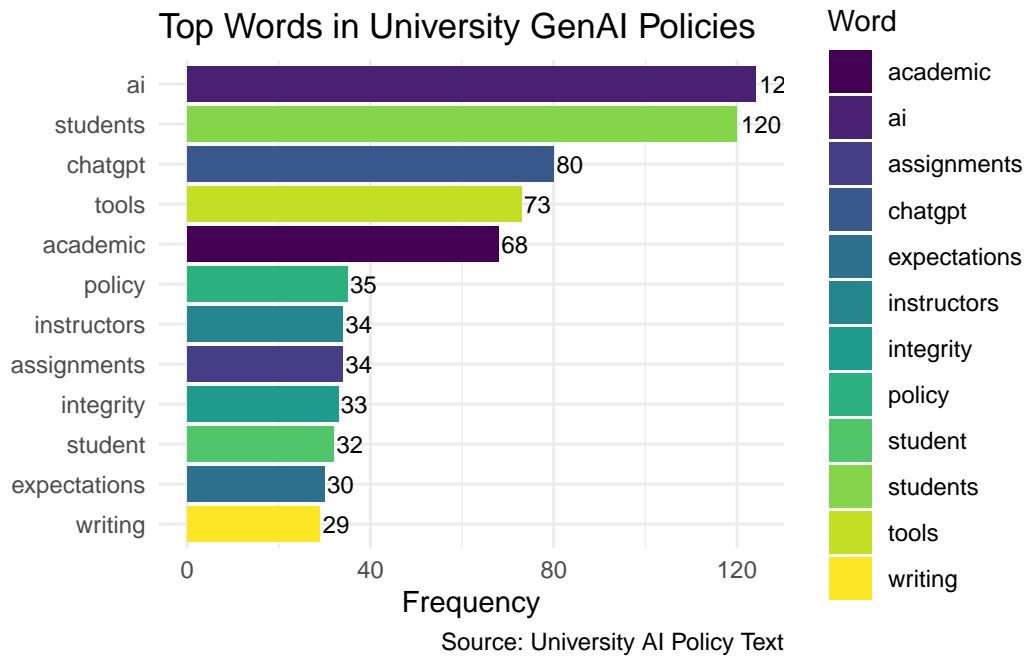
### Step 5: Visualize Top 12 Words in University Policies

Select the top 12 most frequent words and create a bar chart to visualize the distribution.

```
# Select the top 12 words
top_words <- word_frequency %>% slice(1:12)

# Generate the bar chart
policy_word_chart <- ggplot(top_words, aes(x = reorder(word, n), y = n, fill = word)) +
  geom_bar(stat = "identity") +
  coord_flip() +
  theme_minimal() +
  labs(
    title = "Top Words in University GenAI Policies",
    x = NULL,
    y = "Frequency",
    caption = "Source: University AI Policy Text",
    fill = "Word"
  ) +
  scale_fill_viridis(discrete = TRUE) +
  geom_text(aes(label = n), vjust = 0.5, hjust = -0.1, size = 3)

# Print the bar chart
print(policy_word_chart)
```



### 2.3.5 Results and Discussions

- **RQ1:** What are the most frequently mentioned words in university GenAI writing usage policies?

The results of the frequency analysis showed that keywords such as “assignment,” “student,” and “writing” were among the most commonly mentioned terms across AI policies at 100 universities. This emphasis reflects a focus on using AI tools to support student learning and enhance teaching content. The frequent mention of these words suggests that institutions are considering the role of AI in academic assignments and course design, indicating a strategic commitment to integrating AI within educational tasks and student interactions.

- **RQ2:** Which keywords reflect common concerns or focal points related to GenAI writing usage in academic settings?

The analysis of the top 12 frequently mentioned terms highlighted additional focal points, including “tool,” “academic,” “instructor,” “integrity,” and “expectations.” These terms reveal concerns around the ethical use of AI tools, the need for clarity in academic applications, and the central role of instructors in AI policy implementation. Keywords like “integrity” and “expectations” emphasize the importance of maintaining academic standards and setting clear guidelines for AI use in classrooms, while “instructor” underscores the influence faculty members have in shaping AI-related practices. Together, these terms reflect a commitment to transparent policies that support ethical and effective AI integration, enhancing the academic experience for students.

## 2.4 Dictionary-based Analysis

### 2.4.1 Purpose

The purpose of dictionary-based analysis in text data is to assess the presence of predefined categories, like emotions or sentiments, within the text using lexicons or dictionaries. This approach allows researchers to quantify qualitative aspects, such as positive or negative sentiment, based on specific words that correspond to these categories.

**Case:**

In this analysis, we examine the stance of 100 universities on the use of GenAI by applying the Bing sentiment dictionary. By analyzing sentiment scores, we aim to identify the general tone in these policies, indicating whether the institutions' attitudes toward GenAI are predominantly positive or negative.

### 2.4.2 Sample Research Questions

- **RQ:** What is the dominant sentiment expressed in GenAI policy texts across universities, and is it primarily positive or negative?

### 2.4.3 Analysis

#### Step 1: Install and Load Necessary Libraries

First, install and load the required packages for text processing and visualization.

```
# Install necessary packages if not already installed  
#install.packages(c("tidytext", "tidyverse", "dplyr", "ggplot2", "tidyverse"))  
  
# Load libraries  
library(tidytext)  
library(tidyverse)  
library(dplyr)  
library(ggplot2)  
library(tidyverse)
```

## Step 2: Load and Prepare Data(same as 2.3)

Load the GenAI policy stance data from a CSV file. Be sure to update the file path as needed. we use the same data as 2.3.

```
# Load the dataset (replace "University_GenAI_Policy_Stance.csv" with the actual file path)
university_policies <- read_csv("University_GenAI_Policy_Stance.csv")
# Tokenize text, remove stop words, and count word frequencies
word_frequency <- university_policies %>%
  unnest_tokens(word, Stance) %>% # Tokenize the 'Policy_Text' column
  anti_join(stop_words) %>% # Remove common stop words
  count(word, sort = TRUE) # Count and sort words by frequency
word_frequency
```

## Step 3: Tokenize Text Data and Apply Sentiment Dictionary

Tokenize the policy text data to separate individual words. Then, use the Bing sentiment dictionary to label each word as positive or negative.

```
# Tokenize 'Stance' column and apply Bing sentiment dictionary
sentiment_scores <- word_frequency %>%
  inner_join(get_sentiments("bing")) %>% # Join with Bing sentiment lexicon
  pivot_wider(names_from = sentiment, values_from = n, values_fill = 0) %>%
  mutate(sentiment_score = positive - negative) # Calculate net sentiment score

sentiment_scores
```

word	positive	negative	sentiment_score
1 honor	18	0	18
2 cheating	0	11	-11
3 dishonesty	0	10	-10
4 guidance	9	0	9
5 honesty	7	0	7
6 intelligence	7	0	7
7 transparent	7	0	7
8 violation	0	7	-7
9 encourage	6	0	6
10 difficult	0	5	-5
# i 132 more rows			

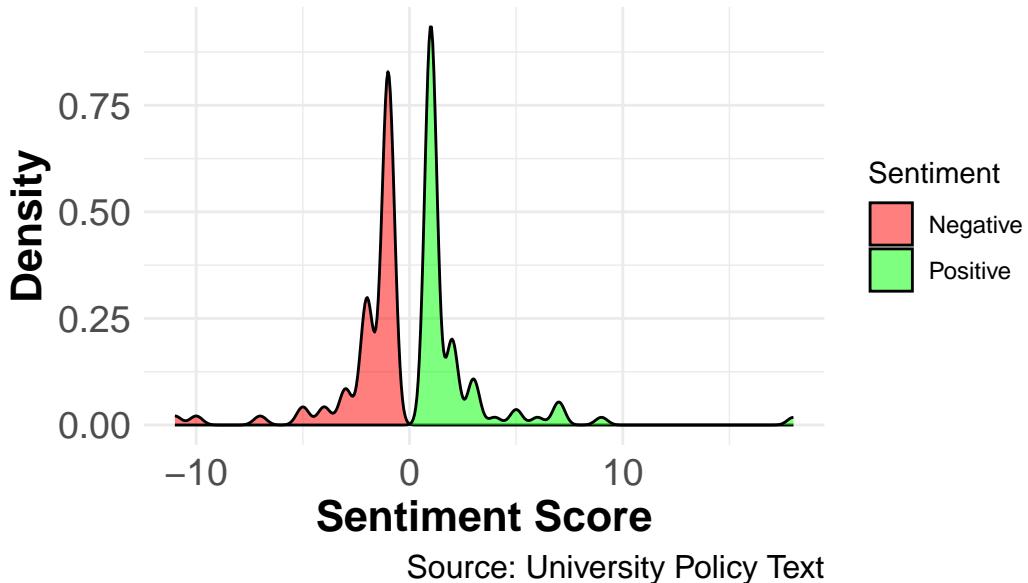
#### Step 4: Create a Density Plot for Sentiment Distribution

Visualize the distribution of sentiment scores with a density plot, showing the prevalence of positive and negative sentiments across university policies.

```
# Generate a density plot of sentiment scores
density_plot <- ggplot(sentiment_scores, aes(x = sentiment_score, fill = sentiment_score > 0))
  geom_density(alpha = 0.5) +
  scale_fill_manual(values = c("red", "green"), name = "Sentiment",
                     labels = c("Negative", "Positive")) +
  labs(
    title = "Density Plot of University AI Policy Sentiment",
    x = "Sentiment Score",
    y = "Density",
    caption = "Source: University Policy Text"
  ) +
  theme_minimal() +
  theme(
    plot.title = element_text(face = "bold", size = 20),
    axis.text = element_text(size = 14),
    axis.title = element_text(face = "bold", size = 16),
    plot.caption = element_text(size = 12)
  )

# Print the plot
print(density_plot)
```

## Density Plot of University AI Policy Sentiment



### 2.4.4 Results and Discussions

- **RQ:** What is the dominant sentiment expressed in GenAI policy texts across universities, and is it primarily positive or negative?

The dictionary-based sentiment analysis reveals the prevailing sentiments in university policies on GenAI usage. Using the Bing lexicon to assign positive and negative scores, the density plot illustrates the distribution of sentiment scores across the 100 institutions.

The results indicate a balanced perspective with a slight tendency toward positive sentiment, as reflected by a higher density of positive scores. This analysis provides insights into the varying degrees of acceptance and caution universities adopt in their AI policy frameworks, demonstrating the diverse stances that shape institutional AI guidelines.

## 2.5 Clustering-Based Analysis

Clustering-based analysis involves grouping similar text documents or text segments into clusters based on their underlying topics or themes. This approach is particularly useful for identifying dominant themes in text data, such as university AI policy documents.

### 2.5.1 Purpose

**Purpose:** The goal of clustering-based analysis is to uncover latent themes in text data using unsupervised machine learning techniques. Topic modeling is one popular method for clustering documents into groups based on their content.

**Case:** Using the GenAI policy texts from 100 universities, we apply Latent Dirichlet Allocation (LDA) to identify dominant themes in these policy documents. This analysis will help categorize policies into overarching themes, such as academic integrity, student support, and instructor discretion.

### 2.5.2 Sample Research Questions

- **RQ1:** What are the prominent themes present in university policies regarding GenAI usage Stance?
- **RQ2:** How do these themes reflect the key concerns or opportunities for integrating GenAI in higher education?

### 2.5.3 Analysis

#### Step 1: Install and Load Necessary Libraries

Install and load the required libraries for text processing and topic modeling.

```
# Install necessary packages
#install.packages(c("dplyr", "tidytext", "topicmodels", "ggplot2"))

# Load libraries
library(dplyr)
library(tidytext)
library(topicmodels)
library(ggplot2)
```

#### Step 2: Prepare the Data

Load the data and create a document-term matrix (DTM) for topic modeling.

```

# Load the dataset
university_policies <- read_csv("University_GenAI_Policy_Stance.csv")

# Tokenize text data and count word frequencies
word_frequency <- university_policies %>%
  unnest_tokens(word, Stance) %>% # same as section 2.3
  anti_join(stop_words) %>% # Remove common stop words
  count(word, sort = TRUE) # Count and sort words by frequency
word_frequency

# A tibble: 1,108 x 2
  word          n
  <chr>     <int>
1 ai            124
2 students      120
3 chatgpt        80
4 tools          73
5 academic       68
6 policy          35
7 assignments    34
8 instructors     34
9 integrity       33
10 student         32
# i 1,098 more rows

# Creating Documents - Word Frequency Matrix
gpt_dtm <- word_frequency %>%
  group_by(word) %>%
  mutate(id = row_number()) %>%
  ungroup() %>%
  cast_dtm(document = "id", term = "word", value = "n")

library(topicmodels)
library(ggplot2)

# Define range of k values
k_values <- c(2, 3, 4, 5)

# Initialize a data frame to store perplexities
perplexities <- data.frame(k = integer(), perplexity = numeric())

```

```

# Calculate perplexity for each k
for (k in k_values) {
  lda_model <- LDA(gpt_dtm, k = k, control = list(seed = 1234)) # Fit LDA model
  perplexity_score <- perplexity(lda_model, gpt_dtm) # Calculate perplexity
  perplexities <- rbind(perplexities, data.frame(k = k, perplexity = perplexity_score))
}

# Plot perplexity vs number of topics
ggplot(perplexities, aes(x = k, y = perplexity)) +
  geom_line() +
  geom_point() +
  labs(
    title = "Perplexity vs Number of Topics",
    x = "Number of Topics (k)",
    y = "Perplexity"
  ) +
  theme_minimal()

```

### Step 3: Fit the LDA Model

Fit an LDA model with  $k = 3$  topics.

```

# Converting document-word frequency matrices to sparse matrices
gpt_dtm_sparse<- as(gpt_dtm, "matrix")

# Fit the LDA model
lda_model <- LDA(gpt_dtm_sparse, k = 3, control = list(seed = 1234))

# View model results
gpt_policy_topics_k3 <- tidy(lda_model, matrix = "beta")

print(gpt_policy_topics_k3)

```

```

# A tibble: 3,324 x 3
  topic term      beta
  <int> <chr>    <dbl>
1     1 ai      0.0408
2     2 ai      0.00593
3     3 ai      0.0650
4     1 students 0.0435
5     2 students 0.0635

```

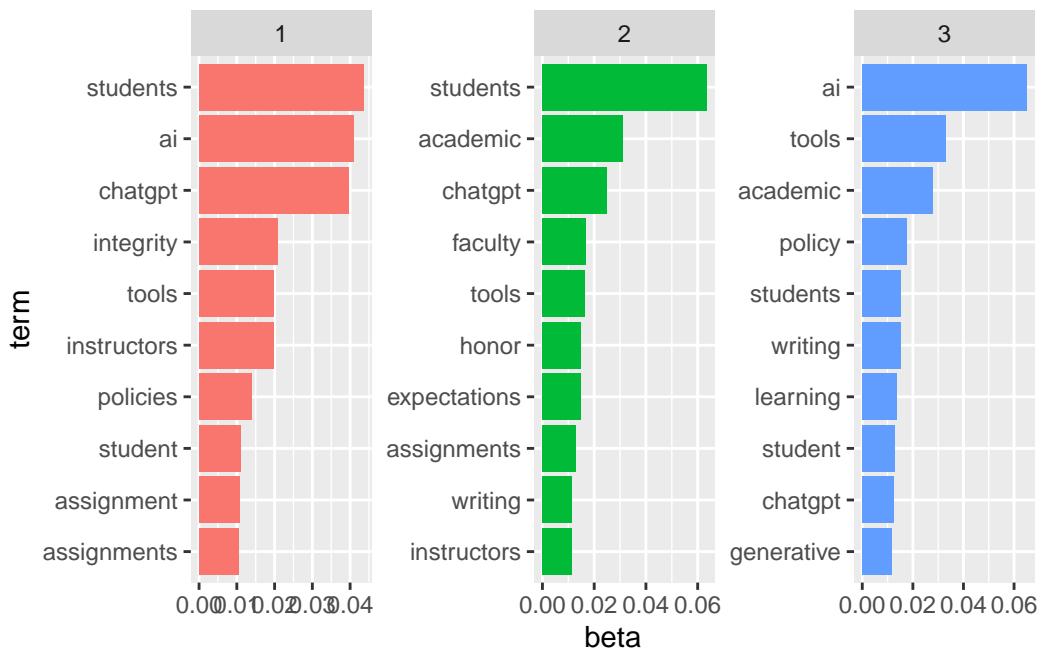
```
6      3 students  0.0151
7      1 chatgpt   0.0396
8      2 chatgpt   0.0248
9      3 chatgpt   0.0126
10     1 tools     0.0197
# i 3,314 more rows
```

#### Step 4: Visualize Topics

Extract the top terms for each topic and visualize them.

```
# Visualizing top terms for each topic
gpt_policy_ap_top_terms_k3 <- gpt_policy_topics_k3 %>%
  group_by(topic) %>%
  slice_max(beta, n = 10) %>%
  ungroup() %>%
  arrange(topic, -beta)

gpt_policy_ap_top_terms_k3 %>%
  mutate(term = reorder_within(term, beta, topic)) %>%
  ggplot(aes(beta, term, fill = factor(topic))) +
  geom_col(show.legend = FALSE) +
  facet_wrap(~ topic, scales = "free") +
  scale_y_reordered()
```



#### 2.5.4 Results and Discussions

##### Research Question 1:

What are the prominent themes present in university policies regarding GenAI usage?

##### Answer:

The topic modeling analysis revealed three distinct themes in the university GenAI policies:

###### 1. Theme 1: Student-Centric Guidelines and Ethical Considerations

- Key terms: **students, integrity, tools, instructors, assignment**
- This theme emphasizes student usage of GenAI in academic settings, with a focus on ethics (**integrity**) and guidelines for instructors to manage assignments involving AI tools.

###### 2. Theme 2: Academic Standards and Faculty Expectations

- Key terms: **students, academic, faculty, honor, expectations**
- This theme focuses on maintaining academic integrity and clarifying expectations for faculty and students regarding GenAI usage in assignments and assessments.

###### 3. Theme 3: Policy-Level Governance and Technology Integration

- Key terms: **ai, tools, policy, learning, generative**
- This theme revolves around institutional policies on AI integration, highlighting broader governance strategies and how generative AI (like GenAI) fits into learning environments.

### **Research Question 2:**

How do these themes reflect the key concerns or opportunities for integrating GenAI in higher education?

#### **Answer:**

The identified themes reflect both concerns and opportunities:

##### **1. Concerns:**

**Theme 1:** Highlights the ethical challenges, such as ensuring academic integrity when students use AI tools in their coursework. Institutions are keen on setting clear guidelines for both students and instructors to avoid misuse.

**Theme 2:** Underlines the potential for conflict between maintaining academic standards (**honor, expectations**) and leveraging AI to support learning. This shows a cautious approach to integrating AI while upholding traditional values.

**Theme 3:** Raises policy-level questions on AI governance, such as whether existing institutional frameworks are adequate to regulate emerging generative AI technologies.

##### **2. Opportunities:**

**Theme 1:** Presents a chance to redefine how students interact with AI tools to foster responsible and innovative usage, particularly for assignments and creative tasks.

**Theme 2:** Encourages collaboration between faculty and administration to develop robust expectations and support systems for integrating AI in the classroom.

**Theme 3:** Offers a strategic opportunity for universities to lead in AI adoption by establishing comprehensive policies that guide AI's role in education and research.

### **Discussion:**

The topic modeling results suggest that universities are navigating a complex landscape of opportunities and challenges as they incorporate GenAI into academic contexts. While student-centric policies aim to balance innovation with ethical considerations, institutional-level themes signal the need for governance frameworks to ensure responsible AI use. These findings indicate that higher education institutions are positioned to play a pivotal role in shaping the future of

generative AI in learning, provided they address the ethical, pedagogical, and policy challenges identified in this analysis.

# Section 3 Social Network Analyses (Relational Data)

## 3.1 Overview

Social network analysis (SNA) is the process of investigating social structures through the use of networks and graph theory. It is a technique used to map and measure relationships and flows between people, groups, organizations, computers, or other information/knowledge processing entities. SNA can be a useful tool for understanding the team structures, for example, in an online classroom. It can be an additional layer of understanding the outcomes (or predictors) of certain instructional interventions. Used this way SNA can be used to identify patterns and trends in social networks, as well as to understand how these networks operate. Additionally, SNA can be used to predict future behavior in social networks, and to design interventions that aim to improve the functioning of these networks.

## 3.2 Accessing SNA Data

Social Network Analysis (SNA) relies on relational data—information about **connections** (edges) between entities (nodes) such as students, teachers, or organizations. Compared to traditional survey or tabular data, SNA requires **pairwise relational information**. In education, this could include “who collaborates with whom,” “who talks to whom,” or digital traces of discussion and collaboration in online platforms.

### 3.2.1 Types and Sources of SNA Data

There are several common sources and structures for SNA data in educational and social science contexts:

- **Survey-based Network Data:** Collected via roster or name generator questions, e.g., “List the classmates you discuss assignments with.”
- **Behavioral/Observational Data:** Derived from logs of actual interactions, e.g., forum replies, emails, classroom seating.
- **Archival or Digital Trace Data:** Extracted from digital platforms such as MOOCs, LMS discussion forums, Slack, Twitter, or Facebook.

- **Administrative/Organizational Data:** Information about formal structures such as team membership or co-authorship.

**Data Structure:** Most SNA data are formatted as: - **Edge List** (two columns: source and target) - **Adjacency Matrix** (rows and columns are actors; cell values indicate a tie) - **Node Attributes** (supplementary information about each node, e.g., gender, role)

### 3.2.2 Example 1: Creating a Simple Network from an Edge List

Below is an example of constructing a network from a simple CSV edge list. This mirrors typical classroom survey data (“who do you consider your friend in this class?”).

```
# Install and load the igraph package
install.packages("igraph")
library(igraph)

# Example: Load an edge list from CSV
edge_list <- read.csv("data/friendship_edges.csv")

# Create the graph object (directed network)
g <- graph_from_data_frame(edge_list, directed = TRUE)

# Plot the network
plot(g, main = "Friendship Network")
```

### 3.2.3 Example 2: Generating Network Data from Digital Traces

Many educational datasets now come from online discussion forums, MOOCs, or LMS systems. For example, the MOOC case study (Kellogg & Edelmann, 2015) uses reply relationships in online courses to construct discussion networks.

```
# Suppose you have a data frame with columns: from_user, to_user
mooc_edges <- read.csv("data/mooc_discussion_edges.csv")
g_mooc <- graph_from_data_frame(mooc_edges, directed = TRUE)
plot(g_mooc, main = "MOOC Discussion Network")
```

### 3.2.4 Example 3: Collecting SNA Data via Surveys

If you want to collect your own network data:

- Ask participants to **name or select** (from a roster) their friends, collaborators, or contacts.
- Compile responses into an edge list.
- Example survey prompt:

“Please list up to five classmates you seek help from most frequently.”

**Tip:**

Survey-based SNA is easier to manage with small to medium groups. For larger networks, digital trace or archival data may be more practical.

### 3.2.5 Node Attribute Data

You can also load additional data about each node (student, teacher, etc.) to enable richer analyses (e.g., centrality by gender or role).

```
node_attributes <- read.csv("data/friendship_nodes.csv")
# Add attributes to igraph object
V(g)$gender <- node_attributes$gender[match(V(g)$name, node_attributes$name)]
```

### 3.2.6 Further Examples

- **Public Datasets:**
  - [MOOC Discussion Networks](#)
  - [Add Health](#)
  - [Common Core Twitter Networks \(Supovitz et al.\)](#)
- **Synthetic Data:**
  - R’s `igraph` package can also generate sample networks for practice:

```
g_sample <- sample_gnp(n = 10, p = 0.3)
plot(g_sample, main = "Random Sample Network")
```

### 3.2.7 Best Practices and Tips

- **Ethics:** Social network data can be sensitive. Protect anonymity and comply with IRB/data use guidelines.
- **Format Consistency:** Always clarify whether ties are directed/undirected, binary/weighted, and ensure consistent formatting.
- **Missing Data:** Especially in survey-based SNA, missing responses can impact network structure and interpretation.

### 3.2.8 Summary

Accessing SNA data involves both careful **design** (in the case of surveys/observations) and **extraction/wrangling** (in the case of digital traces or archival records). The choice of data source and structure will directly influence the kinds of questions you can answer with SNA.

Recommended Reading:

- Borgatti, S. P., Everett, M. G., & Johnson, J. C. (2018). *Analyzing Social Networks* (2nd ed). SAGE.
- Kellogg, S., & Edelmann, A. (2015). Massive open online course discussion forums as networks.

## 3.3 Network Management & Measurement in Social Network Analysis

### 3.3.1 Purpose + Case

**Purpose:** This section demonstrates how to manage, measure, and visualize large-scale discussion networks from online professional development settings. Through this real-world example, we guide readers in loading relational data, constructing a directed network, and conducting a suite of essential SNA measures. The focus is on classroom- and course-level online discussions, which are representative of many contemporary educational and research settings.

**Case Study:** The case data comes from two cohorts of an online professional development program (“DLT1” and “DLT2”). Each cohort’s discussion data includes (a) **edge list** data capturing who replied to whom, and (b) **node/actor** attributes describing roles (e.g., facilitator, expert). These data allow us to reconstruct and analyze the full structure of communication in two authentic online learning communities.

### 3.3.2 Sample Research Questions

- **RQ1:** What is the overall structure of interaction in each online discussion cohort? Are they densely connected, or fragmented?
- **RQ2:** Who are the most central or influential actors in the network? How do facilitators or experts compare with regular participants?
- **RQ3:** To what extent are ties reciprocated (mutual) and how cohesive are the networks?
- **RQ4:** How do the network properties (e.g., density, reciprocity, clustering) compare between cohorts?

### 3.3.3 Analysis

#### Step 1: Install and Load Required Packages

```
# Install and load necessary libraries
#install.packages(c("tidygraph", "ggraph", "readr", "janitor"))
library(tidygraph)
library(ggraph)
library(readr)
library(janitor)
library(igraph)
library(dplyr)
```

#### Step 2: Import and Clean Data Load Edges and Node Attributes for DLT1:

```
# Load edge list (who replied to whom)
dlt1_ties <- read_csv("data/dlt1-edges.csv",
  col_types = cols(Sender = col_character(),
    Receiver = col_character(),
    `Category Text` = col_skip(),
    `Comment ID` = col_character(),
    `Discussion ID` = col_character())) |>
  clean_names()

# Load node attributes (participant roles, etc.)
dlt1_actors <- read_csv("data/dlt1-nodes.csv",
  col_types = cols(UID = col_character(),
    Facilitator = col_character(),
    expert = col_character(),
```

```

    connect = col_character())) |>
clean_names()

head(dlt1_ties)

# A tibble: 6 x 9
  sender receiver timestamp discussion_title discussion_category parent_category
  <chr>   <chr>     <chr>      <chr>          <chr>                  <chr>
1 360     444       4/4/13 1~ Most important ~ Group N           Units 1-3 Disc-
2 356     444       4/4/13 1~ Most important ~ Group D-L        Units 1-3 Disc-
3 356     444       4/4/13 1~ DLT Resources-C~ Group D-L        Units 1-3 Disc-
4 344     444       4/4/13 1~ Most important ~ Group O-T        Units 1-3 Disc-
5 392     444       4/4/13 1~ Most important ~ Group U-Z        Units 1-3 Disc-
6 219     444       4/4/13 1~ Most important ~ Group M           Units 1-3 Disc-
# i 3 more variables: discussion_identifier <chr>, comment_id <chr>,
#   discussion_id <chr>

head(dlt1_actors)

# A tibble: 6 x 13
  uid    facilitator role1 experience experience2 grades location region country
  <chr>  <chr>      <chr>     <dbl> <chr>      <chr>  <chr>  <chr>  <chr>
1 1      0           libme~      1 6 to 10 secon~ VA    South US
2 2      0           class~     1 6 to 10 secon~ FL    South US
3 3      0           distr~     2 11 to 20 gener~ PA    North~ US
4 4      0           class~     2 11 to 20 middle NC   South US
5 5      0           other~     3 20+      gener~ AL    South US
6 6      0           class~     1 4 to 5   gener~ AL    South US
# i 4 more variables: group <chr>, gender <chr>, expert <chr>, connect <chr>
```

### Step 3: Construct and Explore the Network

```

# Build the directed network graph (nodes: uid, edges: sender->receiver)
dlt1_network <- tbl_graph(
  edges = dlt1_ties,
  nodes = dlt1_actors,
  node_key = "uid",
  directed = TRUE
)
```

```

# Overview of the network
dlt1_network

# A tbl_graph: 445 nodes and 2529 edges
#
# A directed multigraph with 4 components
#
# Node Data: 445 x 13 (active)
  uid facilitator role1 experience experience2 grades location region country
  <chr> <chr>      <chr>     <dbl> <chr>      <chr> <chr>      <chr> <chr>
  1 1    0          libm~       1 6 to 10 secon~ VA   South US
  2 2    0          clas~       1 6 to 10 secon~ FL   South US
  3 3    0          dist~       2 11 to 20 gener~ PA   North~ US
  4 4    0          clas~       2 11 to 20 middle NC  South US
  5 5    0          othe~       3 20+      gener~ AL   South US
  6 6    0          clas~       1 4 to 5  gener~ AL   South US
  7 7    0          inst~       2 11 to 20 gener~ SD   Midwe~ US
  8 8    0          spec~       1 6 to 10 secon~ BE   Inter~ BE
  9 9    0          clas~       1 6 to 10 middle NC  South US
 10 10   0          scho~       2 11 to 20 middle NC  South US
# i 435 more rows
# i 4 more variables: group <chr>, gender <chr>, expert <chr>, connect <chr>
#
# Edge Data: 2,529 x 9
  from    to timestamp discussion_title discussion_category parent_category
  <int> <int> <chr>      <chr>           <chr>           <chr>
  1 360  444 4/4/13 16:32 Most important c~ Group N           Units 1-3 Disc~
  2 356  444 4/4/13 18:45 Most important c~ Group D-L         Units 1-3 Disc~
  3 356  444 4/4/13 18:47 DLT Resources-Co~ Group D-L         Units 1-3 Disc~
# i 2,526 more rows
# i 3 more variables: discussion_identifier <chr>, comment_id <chr>,
#   discussion_id <chr>

# Output: 445 nodes, 2529 edges, 4 components (directed multigraph)


```

#### Step 4: Basic Visualization


```

# Quick overview plot (stress layout by default)
autograph(dlt1_network)

```



```
# Custom visualization with colors and centrality
graph(dlt1_network, layout = "fr") +
  geom_edge_link(alpha = .2) +
  geom_node_point(aes(color = role1, size = local_size())) +
  theme_graph()+
  theme(text = element_text(family = "sans"))
```



## Step 5: Network Size and Centralization

```
# Number of nodes and edges  
gorder(dlt1_network) # 445  
  
[1] 445  
  
gsizes(dlt1_network) # 2529  
  
[1] 2529  
  
# Degree centrality (all, in, out)  
deg_all <- centr_degree(dlt1_network, mode = "all")$res  
deg_in <- centr_degree(dlt1_network, mode = "in")$res  
deg_out <- centr_degree(dlt1_network, mode = "out")$res  
# Centralization  
centr_degree(dlt1_network, mode = "all")$centralization # 0.64
```

```
[1] 0.6429242
```

```
centr_degree(dlt1_network, mode = "in")$centralization # 1.06
```

```
[1] 1.05702
```

```
centr_degree(dlt1_network, mode = "out")$centralization # 0.23
```

```
[1] 0.2259389
```

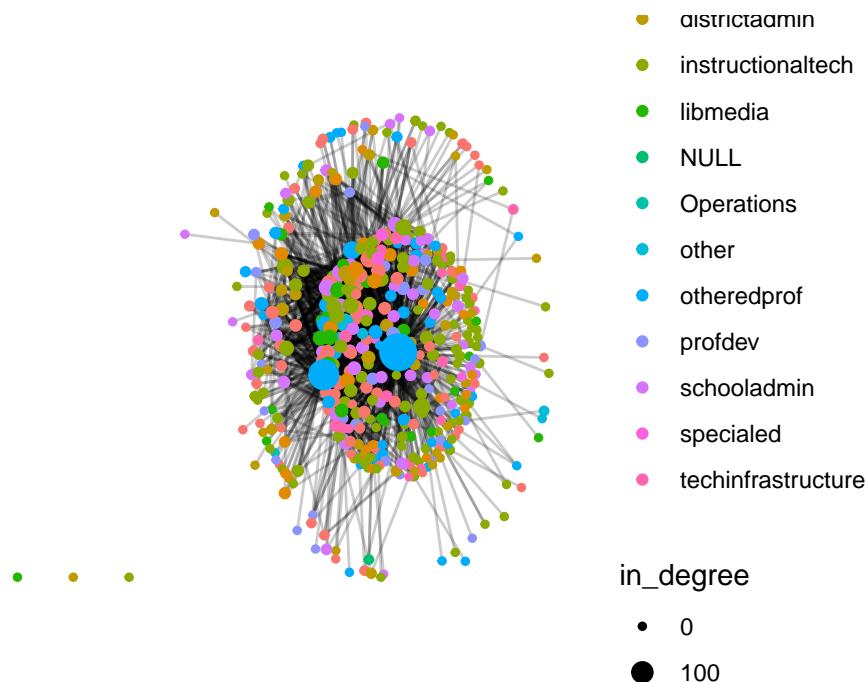
## Step 6: Attach and Visualize Node Centrality

```
# Add in-degree centrality as node attribute  
dlt1_network <- dlt1_network |>  
  activate(nodes) |>  
  mutate(in_degree = centrality_degree(mode = "in"))  
  
# Plot, sizing nodes by in-degree
```

```

ggraph(dlt1_network) +
  geom_edge_link(alpha = .2) +
  geom_node_point(aes(size = in_degree, color = role1)) +
  theme_graph()+
  theme(text = element_text(family = "sans"))

```



### Step 7: Network Density, Reciprocity, Clustering, Distance

```

# Density
edge_density(dlt1_network)      # 0.013 (sparse network)

```

```
[1] 0.01279988
```


```

# Reciprocity
reciprocity(dlt1_network)      # 0.20 (20% of ties are reciprocated)

```

```
[1] 0.1997544
```

```
# Add reciprocated edge attribute and plot
dlt1_network <- dlt1_network |>
  activate(edges) |>
  mutate(reciprocated = edge_is_mutual())
ggraph(dlt1_network) +
  geom_node_point(aes(size = in_degree)) +
  geom_edge_link(aes(color = reciprocated), alpha = .2) +
  theme_graph()+
  theme(text = element_text(family = "sans"))
```



```
# Clustering (transitivity/global)
transitivity(dlt1_network)      # 0.089
```

```
[1] 0.08880774
```

```
# Network diameter (longest shortest path) & average distance
diameter(dlt1_network)          # 8
```

```
[1] 8
```

```
mean_distance(dlt1_network)      # 3.03
```

```
[1] 3.030694
```

## Step 8:Repeat for DLT2

```
# Step 8: Repeat for DLT2

# 1. Load the DLT2 edge and node data
dlt2_ties <- read_csv("data/dlt2-edges.csv",
  col_types = cols(Sender = col_character(),
                    Receiver = col_character(),
                    `Category Text` = col_skip(),
                    `Comment ID` = col_character(),
                    `Discussion ID` = col_character())) |>
  clean_names()

dlt2_actors <- read_csv("data/dlt2-nodes.csv",
  col_types = cols(UID = col_character(),
                    Facilitator = col_character(),
                    expert = col_character(),
                    connect = col_character())) |>
  clean_names()

# 2. Construct the directed network
dlt2_network <- tbl_graph(
  edges = dlt2_ties,
  nodes = dlt2_actors,
  node_key = "uid",
  directed = TRUE
)

# 3. Basic network properties
num_nodes <- gorder(dlt2_network)    # Number of nodes
num_edges <- gsize(dlt2_network)      # Number of edges

# 4. Degree centrality (overall, in, out)
deg_all <- centr_degree(dlt2_network, mode = "all")$res
deg_in <- centr_degree(dlt2_network, mode = "in")$res
deg_out <- centr_degree(dlt2_network, mode = "out")$res


# Centralization values
centr_all <- centr_degree(dlt2_network, mode = "all")$centralization
centr_in <- centr_degree(dlt2_network, mode = "in")$centralization
centr_out <- centr_degree(dlt2_network, mode = "out")$centralization
```

```

# 5. Attach centrality as a node attribute
dlt2_network <- dlt2_network |>
  activate(nodes) |>
  mutate(in_degree = centrality_degree(mode = "in"))

# 6. Visualize the network
ggraph(dlt2_network, layout = "fr") +
  geom_edge_link(alpha = .2) +
  geom_node_point(aes(color = role, size = local_size())) +
  theme_graph() +
  theme(text = element_text(family = "sans"))

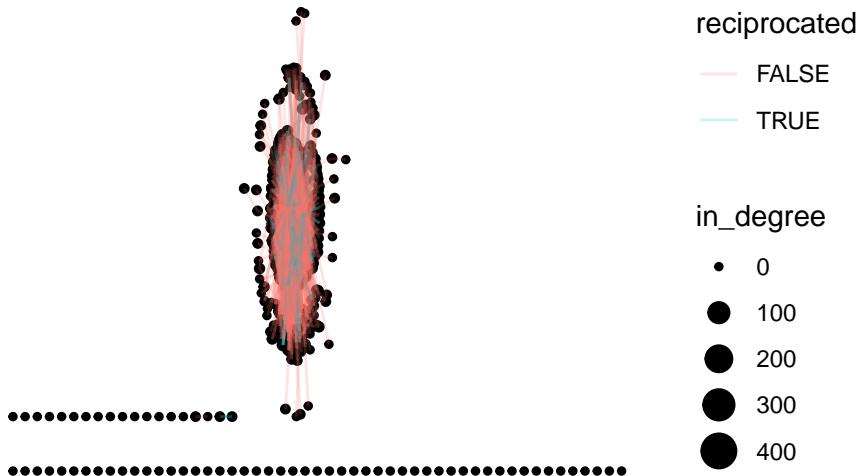
```



```

# 7. Density, reciprocity, clustering, distances
density <- edge_density(dlt2_network)
recip   <- reciprocity(dlt2_network)
dlt2_network <- dlt2_network |>
  activate(edges) |>
  mutate(reciprocated = edge_is_mutual())
ggraph(dlt2_network) +
  geom_node_point(aes(size = in_degree)) +
  geom_edge_link(aes(color = reciprocated), alpha = .2) +
  theme_graph() +
  theme(text = element_text(family = "sans"))

```



```

trans   <- transitivity(dlt2_network)
diam    <- diameter(dlt2_network)
mean_d  <- mean_distance(dlt2_network)

# 8. Print summary statistics
cat("DLT2 Network Stats:\n")

```

DLT2 Network Stats:

```
cat("Nodes:", num_nodes, "Edges:", num_edges, "\n")
```

Nodes: 492 Edges: 2584

```
cat("Degree Centralization (all/in/out):", centr_all, centr_in, centr_out, "\n")
```

Degree Centralization (all/in/out): 0.5311161 0.8650671 0.3273889

```
cat("Density:", density, "Reciprocity:", recip, "\n")
```

Density: 0.0106966 Reciprocity: 0.2500977

```
cat("Transitivity:", trans, "Diameter:", diam, "Mean Distance:", mean_d, "\n")
```

Transitivity: 0.1248291 Diameter: 8 Mean Distance: 3.03815

### 3.3.4 Results and Discussion

#### RQ1: What is the overall structure of interaction in each cohort?

- DLT1 consists of **445 nodes** (participants) and **2529 edges** (directed interactions).
- DLT2 has **492 nodes** and **2584 edges**.
- Both networks are **large and sparse**:
  - *Density*: DLT1 = **0.013**, DLT2 = **0.011**
  - Interpretation: Only about 1–1.3% of all possible connections exist—typical for online discussion networks where not every participant interacts with every other.
- Both networks are **multi-component** (several disconnected groups), but most participants are included in the main giant component.
- The **diameter** (longest shortest path) is 8 for both cohorts, and the **average shortest path length** is about **3.03 (DLT1)** and **3.04 (DLT2)**, indicating that on average, any participant is just 3 steps away from any other in the largest group.
- **Interpretation:** Information or discussion threads can reach most participants with only a few hops, but overall engagement is selective rather than comprehensive.

#### RQ2: Who are the most central or influential actors?

- Centrality (degree, in-degree, out-degree) analyses show a **right-skewed distribution**: most participants have low centrality, but a small subset are highly connected.
- In both DLT1 and DLT2, **facilitators** and a handful of highly active participants emerge as **hubs**—they initiate and/or receive a disproportionate number of interactions.
  - For DLT1, degree centralization (all): **0.64** (in: **1.06**, out: **0.23**)
  - For DLT2, degree centralization (all): **0.53** (in: **0.87**, out: **0.33**)
- **Visualization:** Network plots with node size proportional to in-degree clearly highlight these central actors.
- **Interpretation:** These key individuals (often facilitators) play critical roles in steering discussion, providing feedback, and potentially keeping less active members engaged.

#### RQ3: Are ties reciprocated?

- **Reciprocity** (proportion of mutual connections):
  - DLT1: **0.20** (20% of ties are reciprocated)
  - DLT2: **0.25** (25% reciprocated)

- **Interpretation:** Most interactions are one-way (e.g., a reply that does not receive a response), but a substantial fraction are mutual—possibly reflecting peer-to-peer conversations or ongoing exchanges. In online learning contexts, this suggests a mix of broadcasting (one-to-many) and genuine dialog (two-way).

#### RQ4: How cohesive are the networks?

- **Transitivity/Clustering coefficient** (probability that two connected nodes' neighbors are also connected):
  - DLT1: **0.089**
  - DLT2: **0.125**
- **Interpretation:** Triadic closure is low—there are few closed triangles, so close-knit groups (where “my friend is also your friend”) are rare. The network structure is more “hub-and-spoke” than “cliquish.”
- **Diameter:** 8 for both, showing that even the furthest nodes can be reached in 8 steps.
- **Mean distance:** ~3.0, so participants are relatively close to each other in the main component.

#### Comparison Across Cohorts

- **DLT2 is slightly larger** (more participants and interactions), but the structural properties—density, centralization, reciprocity, clustering, and path lengths—are all quite similar.
- Minor variations (e.g., higher reciprocity and clustering in DLT2) could reflect differences in facilitation style, cohort engagement, or participant composition.
- **Interpretation:** Both cohorts exhibit classic patterns for large-scale online educational discussions—a small number of central actors drive most of the interaction, the networks are sparse but efficiently connected, and genuine dialogue is present but not universal.

#### Educational Implications

- **For educators and instructional designers:**

These findings suggest that a small group of highly active facilitators or students are critical to fostering interaction. Encouraging more distributed engagement (for example, through peer response requirements or rotating leadership) may enhance network cohesion and learning opportunities.

- **For researchers:**

Understanding who occupies central positions and the overall structure of discussion networks can inform interventions to support isolated participants, promote reciprocity, and create more connected learning communities.

## **Summary:**

Through these SNA measures, we have shown how to reconstruct, visualize, and interpret the structure of large-scale online discussion networks. The approach enables identification of core communicators, understanding of participation patterns, and empirical comparison across cohorts or interventions. This “cookbook” can be adapted to other online learning or collaborative contexts.

> **Note:** This analysis is based on real-world data from online professional development courses. The methods and findings can be generalized to other educational settings where social networks play a role in learning and collaboration.

## **3.4 Case Study: Hashtag Common Core**

### **3.4.1 Purpose & Case**

The purpose of this case study is to demonstrate the application of social network analysis (SNA) in a real-world policy context: the heated national debate over the Common Core State Standards (CCSS) as it played out on Twitter. Drawing on the work of Supovitz, Daly, del Fresno, and Kolouch, the #COMMONCORE Project provides a vivid example of how social media-enabled networks shape educational discourse and policy.

This case focuses on: - Identifying **key actors** (“transmitters,” “transceivers,” and “transcenders”) and measuring their influence, - Detecting **subgroups/factions** within the conversation, - Exploring how **sentiment** about the Common Core varies across network positions, - Demonstrating **network wrangling, visualization, and analysis** using real tweet data.

### **Data Source**

Data was collected from Twitter’s public API using keywords/hashtags related to the Common Core (e.g., **#commoncore**, **ccss**, **stopcommoncore**). The dataset includes user names, tweets, mentions, retweets, and relevant timestamps from a sample week. Only public tweets are included, and user privacy is respected.

### **3.4.2 Sample Research Questions**

- **RQ1:** Who are the “transmitters,” “transceivers,” and “transcenders” in the Common Core Twitter network?
- **RQ2:** What subgroups or factions exist within the network, and how are they structured?
- **RQ3:** How does sentiment about the Common Core vary across actors and subgroups?

- **RQ4:** What other patterns of communication (e.g., centrality, clique formation, isolates) characterize this network?

### 3.4.3 Analysis

#### Step 1: Load Required Packages

```
library(tidyverse)
library(tidygraph)
library(ggraph)
library(skimr)
library(igraph)
library(tidytext)
library(vader)
```

#### Step 2: Data Import and Wrangling

```
# Import tweet data (edgelist format: sender, receiver, timestamp, text)
ccss_tweets <- read_csv("data/ccss-tweets.csv")

# Prepare the edgelist (extract sender, mentioned users, and tweet text)
ties_1 <- ccss_tweets %>%
  relocate(sender = screen_name, target = mentions_screen_name) %>%
  select(sender, target, created_at, text)

# Unnest receiver to handle multiple mentions per tweet
ties_2 <- ties_1 %>%
  unnest_tokens(input = target,
                output = receiver,
                to_lower = FALSE) %>%
  relocate(sender, receiver)

# Remove tweets without mentions to focus on direct connections
ties <- ties_2 %>%
  drop_na(receiver)

# Save for reproducibility
write_csv(ties, "data/ccss-edgelist.csv")
```

```

# Build nodelist
actors_1 <- ties %>%
  select(sender, receiver) %>%
  pivot_longer(cols = c(sender,receiver))

actors <- actors_1 %>%
  select(value) %>%
  rename(actors = value) %>%
  distinct()

write_csv(actors, "data/ccss-nodelist.csv")

```

### Step 3: Create Network Object

```

ccss_network <- tbl_graph(edges = ties,
                           nodes = actors,
                           directed = TRUE)
ccss_network

```

```

# A tbl_graph: 46 nodes and 42 edges
#
# A directed multigraph with 14 components
#
# Node Data: 46 x 1 (active)
  actors
  <chr>
1 DistanceLrnBot
2 k12movieguides
3 WEquilSchool
4 JoeWEquil
5 SumayLu
6 fluttbott
7 BodShameless
8 Math
9 ozsultan
10 sfchronicle
# i 36 more rows
#
# Edge Data: 42 x 4
  from      to created_at          text

```

```

<int> <int> <dttm>           <chr>
1     1     2 2021-06-28 09:53:54 "#Luca Movie Guide | Worksheet | Questions | ~
2     3     4 2021-06-28 02:32:59 "Why public schools should focus more on buil~
3     3     3 2021-06-28 02:32:59 "Why public schools should focus more on buil~
# i 39 more rows

```

#### Step 4: Network Structure – Components, Cliques, and Communities

- Components

- Identify weak and strong components (connected subgroups):

```

ccss_network <- ccss_network |>
  activate(nodes) |>
  mutate(weak_component = group_components(type = "weak"),
         strong_component = group_components(type = "strong"))
# View component sizes
ccss_network |>
  as_tibble() |>
  group_by(weak_component) |>
  summarise(size = n()) |>
  arrange(desc(size))

# A tibble: 14 x 2
  weak_component size
  <int> <int>
1 1     14
2 2     6
3 3     4
4 4     3
5 5     3
6 6     2
7 7     2
8 8     2
9 9     2
10 10    2
11 11    2
12 12    2
13 13    1
14 14    1

```

- Cliques

- Identify fully connected subgroups (if any):

```
clique_num(ccss_network)
```

```
[1] 4
```

```
cliques(ccss_network, min = 3)
```

```
[[1]]  
+ 3/46 vertices, from 29a301e:
```

```
[1] 4 5 6
```

```
[[2]]
```

```
+ 3/46 vertices, from 29a301e:
```

```
[1] 39 40 41
```

```
[[3]]
```

```
+ 3/46 vertices, from 29a301e:
```

```
[1] 3 4 6
```

```
[[4]]
```

```
+ 4/46 vertices, from 29a301e:
```

```
[1] 3 4 5 6
```

```
[[5]]
```

```
+ 3/46 vertices, from 29a301e:
```

```
[1] 3 4 5
```

```
[[6]]
```

```
+ 3/46 vertices, from 29a301e:
```

```
[1] 3 5 6
```

- **Communities**

- Detect densely connected communities using edge betweenness:

```
ccss_network <- ccss_network |>  
  morph(to_undirected) |>  
  activate(nodes) |>  
  mutate(sub_group = group_edge_betweenness()) |>  
  unmorph()
```

```

ccss_network |>
  as_tibble() |>
  group_by(sub_group) |>
  summarise(size = n()) |>
  arrange(desc(size))

```

```

# A tibble: 16 x 2
  sub_group size
  <int>     <int>
1       1     10
2       2      6
3       3      4
4       4      3
5       5      3
6       6      2
7       7      2
8       8      2
9       9      2
10      10     2
11      11     2
12      12     2
13      13     2
14      14     2
15      15     1
16      16     1

```

### Step 5: Egocentric Analysis – Centrality & Key Actors

```

ccss_network <- ccss_network |>
  activate(nodes) |>
  mutate(
    size = local_size(),
    in_degree = centrality_degree(mode = "in"),
    out_degree = centrality_degree(mode = "out"),
    closeness = centrality_closeness(),
    betweenness = centrality_betweenness()
  )

# Identify top actors by out_degree (transmitters), in_degree (transceivers), and both (trans
top_transmitters <- ccss_network %>% as_tibble() %>% arrange(desc(out_degree)) %>% head(5)

```

```

top_transceivers <- ccss_network %>% as_tibble() %>% arrange(desc(in_degree)) %>% head(5)
top_transcenders <- ccss_network %>% as_tibble() %>%
  filter(out_degree > quantile(out_degree, 0.9) & in_degree > quantile(in_degree, 0.9))


```

### Step 6: Visualize the Network

```

ggraph(ccss_network, layout = "fr") +
  geom_node_point(aes(size = out_degree, color = out_degree)) +
  geom_edge_link(alpha = .2) +
  theme_graph()+
  theme(text = element_text(family = "sans"))

```



### Step 7: Sentiment Analysis (Optional)

If you want to analyze sentiment as in the original #COMMONCORE study:

```

library(vader)
vader_ccss <- vader_df(ccss_tweets$text)
mean(vader_ccss$compound)

```

[1] 0.08668182

```

vader_ccss_summary <- vader_ccss %>%
  mutate(sentiment = case_when(
    compound >= 0.05 ~ "positive",
    compound <= -0.05 ~ "negative",
    TRUE ~ "neutral"
  )) %>%
  count(sentiment)

```

### 3.4.4 Results and Discussion

#### RQ1: Who are the “transmitters,” “transceivers,” and “transcenders” in the Common Core Twitter network?

- **Transmitters (high out-degree):**

The user SumayLu stands out as the top transmitter, initiating 8 outgoing ties (mentions/retweets), followed by DouglasHolt... (5), WEquilSchool (3), fluttbot (3), and JoeWEquil (2). These users are the most active in broadcasting or mentioning others within the network.

- **Transceivers (high in-degree):**

The most-mentioned users are WEquilSchool and SumayLu (in-degree = 3), JoeWEquil (2), Tech4Learni... (2), and LASER\_Insti... (2). These individuals receive the most attention from other actors—potential focal points in conversations.

- **Transcenders (high in-degree and out-degree):**

Only two users—WEquilSchool (in-degree = 3, out-degree = 3) and SumayLu (in-degree = 3, out-degree = 8)—simultaneously act as hubs for both sending and receiving communication. These “bridging” actors may serve as key facilitators or connectors in the discourse.

#### RQ2: What subgroups or factions exist in the network?

- **Component analysis** shows a **fragmented network**:

- There are **14 weakly connected components**, the largest containing 14 users, and several small groups or dyads (many with just 2–3 members).
- This fragmentation suggests limited overall cohesion, with multiple parallel or isolated conversations occurring.

- **Clique analysis** reveals:

- Four cliques (fully connected subgroups) of size 3 or 4—e.g., one 4-person clique involving nodes 3, 4, 5, and 6, and several overlapping 3-person cliques. This indicates pockets of tight-knit interaction, but such groups are rare relative to the size of the network.
- **Community detection** using edge betweenness identifies 16 subgroups, generally aligning with the component structure. The largest subgroup has 10 members, with most others much smaller.

### RQ3: What is the overall sentiment in the network?

- **VADER sentiment analysis** of tweet content yields:
  - An average sentiment score (**compound**) of **0.09** (slightly positive), indicating that, despite the policy controversy, the sampled tweets were, on balance, more positive than negative.
  - When tweets are classified into categories:
    - \* A mix of *positive*, *neutral*, and *negative* tweets is observed, with positive tweets slightly outnumbering negatives.
  - This suggests the debate, at least in this time slice, included advocacy and constructive dialogue, not only criticism or negativity.

### RQ4: What other patterns of communication (e.g., centrality, clique formation, isolates) characterize this network?

- **Centrality Patterns:**

The network displays a classic “star” structure in its largest component. Two users, SumayLu and WEquilSchool, stand out with high out-degree and in-degree centrality, respectively. Most other users have very low degree values (often 0 or 1), meaning they are peripheral, engaging in few interactions.

- *Transmitters* (high out-degree): e.g., SumayLu (8 outgoing ties), DouglasHolt... (5).
- *Transceivers* (high in-degree): e.g., WEquilSchool, SumayLu (both in-degree = 3).
- *Transcenders* (both high in- and out-degree): rare—only WEquilSchool and SumayLu meet this criterion in this sample.

- **Clique Formation:**

Clique analysis revealed **4 cliques** (fully connected subgroups) of size 3 or more, with one larger clique (size 4) and several overlapping smaller cliques. However, cliques are rare and limited in size—most communication occurs outside of dense subgroups.

- **Isolates and Components:**

The network has **14 weak components**—many of them tiny. Several users are **isolates** or part of isolated dyads and triads, meaning they are disconnected from the main conversation or only loosely connected. This points to a lack of broad, network-wide cohesion.

- **Community Structure:**

Edge betweenness community detection found 16 subgroups, typically matching up with the component structure: most subgroups are very small (2–3 nodes), while the largest subgroup consists of 10 users.

- **Summary:**

Communication in this network is characterized by:

- **Strong centralization** around a small number of users (hubs);
- **Sparse and fragmented structure** with many small, disconnected components;
- **Limited clique formation**—pockets of tightly connected users exist but are rare;
- **Numerous isolates**—users who are only weakly or not at all connected to the core discussion.

## Discussion

This analysis of the Common Core Twitter conversation reveals a **sparse and fragmented network structure**. The debate is distributed across many small subgroups, with only one moderately sized component (14 members). Within this landscape:

- **Key actors** such as SumayLu and WEquilSchool serve as both broadcasters and focal points of attention (“transcenders”), but most users are peripheral, interacting minimally.
- **Cliques and communities** are few and small, underscoring the lack of broad cohesion. Most interactions happen within micro-groups rather than across the entire network.
- **Sentiment** is, perhaps surprisingly, slightly positive on average. This may reflect the presence of advocacy groups, promotional messaging, or simply a lack of highly negative engagement during the observed period.

## Implications:

The findings illustrate classic social network phenomena in online policy debate: - Most users are only lightly involved, and only a select few drive discussion or receive significant attention. - Communication is siloed, with many small isolated groups and minimal bridging between them. - Sentiment analysis offers nuance: while public debates may be assumed to be contentious, the prevailing tone can still be balanced or even positive in certain time slices.

**For researchers and practitioners**, this means that: - Identifying and engaging “transcenders” is essential for bridging subgroups and spreading information. - Interventions or

outreach should consider the network's fragmentation—broader influence may require engaging multiple small groups individually rather than targeting a single “core.” - Combining SNA with text/sentiment analysis gives a fuller picture: not just who is talking, but *how* and *with what tone*.

Future analysis could track changes in sentiment and connectivity over time, or compare subgroups for differences in message tone and network position.

## References

- Supovitz, J., Daly, A.J., del Fresno, M., & Kolouch, C. (2017). *#commoncore Project*. Retrieved from <http://www.hashtagcommoncore.com>
- Carolan, B.V. (2014). *Social Network Analysis and Education: Theory, Methods & Applications*. Sage.
- Silge, J., & Robinson, D. (2017). *Text Mining with R: A Tidy Approach*. O'Reilly.

# **Section 4 Secondary Analysis of Big Data (Numeric Data)**

**Abstract:** This section reviews how to access data that is primarily numeric/quantitative in nature, but from a different source and of a different nature than the data typically used by social scientists. Example data sets include international or national large-scale assessments (e.g., PISA, NAEP IPEDS) and data from digital technologies (e.g., log-trace data from Open University Learning Analytics Dataset (OULAD)).

## **4.1 Overview**

In social science research, data is traditionally sourced from small-scale surveys, experiments, or qualitative studies. However, the rise of big data offers researchers opportunities to explore numeric and quantitative datasets of unprecedented scale and variety. This chapter discusses how to access and analyze large-scale datasets like international assessments (e.g., PISA, NAEP) and digital log-trace data (e.g., Open University Learning Analytics Dataset (OULAD)). These secondary data sources enable novel research questions and methods, particularly when paired with machine learning and statistical modeling approaches.

## **4.2 Accessing Big data (Broadening the Horizon)**

### **4.2.1 Big Data**

#### **Accessing PISA Data**

The Programme for International Student Assessment (PISA) is a widely used dataset for large-scale educational research. It assesses 15-year-old students' knowledge and skills in reading, mathematics, and science across multiple countries. Researchers can access PISA data through various methods:

## **1. Direct Download from the Official Website**

The OECD provides direct access to PISA data files via its official website. Researchers can download data for specific years and cycles. Data files are typically provided in .csv or .sav (SPSS) formats, along with detailed documentation.

- **Steps to Access PISA Data from the OECD Website:**

1. Visit the [OECD PISA website](#).
2. Navigate to the “Data” section.
3. Select the desired assessment year (e.g., 2022).
4. Download the data and accompanying codebooks.

## **2. Using the OECD R Package**

The OECD R package provides a direct interface to download and explore datasets published by the OECD, including PISA.

- **Steps to Use the OECD Package:**

1. Install and load the OECD package.
2. Use the `getOECD()` function to fetch PISA data.

```
# Install and load the OECD package
install.packages("OECD")
library(OECD)

# Fetch PISA data for the 2018 cycle
pisa_data <- getOECD("pisa", years = "2022")

# Display a summary of the data
summary(pisa_data)
```

## **3. Using the Edsurvey R Package**

The Edsurvey package is designed specifically for analyzing large-scale assessment data, including PISA. It allows for complex statistical modeling and supports handling weights and replicate weights used in PISA.

- **Steps to Use the Edsurvey Package:**

1. Install and load the Edsurvey package.
2. Download the PISA data from the OECD website and provide the path to the .sav files.
3. Load the data into R using `readPISA()`.

```

# Install and load the Edsurvey package
install.packages("Edsurvey")
library(Edsurvey)

# Read PISA data from a local file
pisa_data <- readPISA("path/to/PISA2022Student.sav")

# Display the structure of the dataset
str(pisa_data)

```

## Comparison of Methods

Method	Advantages	Disadvantages
Direct Download	Full access to all raw data and documentation.	Requires manual processing and cleaning.
OECD Package	Easy to use for downloading specific datasets.	Limited to OECD-published formats.
Edsurvey Package	Supports advanced statistical analysis and weights.	Requires additional setup and dependencies.

## Accessing IPEDS Data

The Integrated Postsecondary Education Data System (IPEDS) is a comprehensive source of data on U.S. colleges, universities, and technical and vocational institutions. It provides data on enrollments, completions, graduation rates, faculty, finances, and more. Researchers and policymakers widely use IPEDS data to analyze trends in higher education.

There are several ways to access IPEDS data, depending on the user's needs and technical proficiency.

### 1. Direct Download from the NCES Website

The most straightforward way to access IPEDS data is by downloading it directly from the National Center for Education Statistics (NCES) website.

#### Steps to Access IPEDS Data:

1. Visit the [IPEDS Data Center](#).
2. Click on “Use the Data” and navigate to the “Download IPEDS Data Files” section.

3. Select the desired data year and survey component (e.g., Fall Enrollment, Graduation Rates).
4. Download the data files, typically provided in .csv or .xls format, along with accompanying codebooks.

## **2. Using the ipeds R Package**

The `ipeds` R package simplifies downloading and analyzing IPEDS data directly from R by connecting to the NCES data repository.

### **Steps to Use the ipeds Package:**

1. Install and load the `ipeds` package.
2. Use the `download_ipeds()` function to fetch data for specific survey components and years.

```
# Install and load the ipeds package
install.packages("ipeds")
library(ipeds)

# Download IPEDS data for completions in 2021
ipeds_data <- download_ipeds("C", year = 2021)

# View the structure of the downloaded data
str(ipeds_data)
```

## **3. Using the tidycensus R Package**

The `tidycensus` package, while primarily designed for Census data, can access specific IPEDS data linked to educational institutions.

### **Steps to Use the tidycensus Package:**

1. Install and load the `tidycensus` package.
2. Set up a Census API key to access the data.
3. Query IPEDS data for specific institution-level information.

```

# Install and load the tidycensus package
install.packages("tidycensus")
library(tidycensus)

# Set Census API key (replace with your actual key)
census_api_key("your_census_api_key")

# Fetch IPEDS-related data (e.g., institution information)
ipeds_institutions <- get_acs(
  geography = "place",
  variables = "B14002_003",
  year = 2021,
  survey = "acs5"
)

# View the first few rows
head(ipeds_institutions)

```

#### **4. Using Online Tools**

IPEDS provides several online tools for querying and visualizing data without requiring programming skills.

##### **Common Tools:**

- **IPEDS Data Explorer:** Enables users to query and export customized datasets.
- **Trend Generator:** Allows users to visualize trends in key metrics over time.
- **IPEDS Use the Data:** Simplified tool for accessing pre-compiled datasets.

##### **Steps to Use the IPEDS Data Explorer:**

1. Visit the IPEDS Data Explorer.
2. Select variables of interest, such as institution type, enrollment size, or location.
3. Filter results by years, institution categories, or other criteria.
4. Export the results as a .csv or .xlsx file.

##### **Comparison of Methods**

Method	Advantages	Disadvantages
Direct Download	Full access to raw data and documentation.	Requires manual data preparation and cleaning.
<code>ipeds</code> Package	Automated access to specific components.	Limited flexibility for customized queries.
<code>tidycensus</code> Package	Allows integration with Census and ACS data.	Requires API setup and advanced R skills.
Online Tools	User-friendly and suitable for non-coders.	Limited to predefined queries and exports.

## Accessing Open University Learning Analytics Dataset (OULAD)

The **Open University Learning Analytics Dataset (OULAD)** is a publicly available dataset designed to support research in educational data mining and learning analytics. It includes student demographics, module information, interactions with the virtual learning environment (VLE), and assessment scores.

### Steps to Access OULAD Data

#### Visit the OULAD Repository\*\*

The dataset is hosted on the [Open University's Analytics Project](#). To access the data: 1. Navigate to the website. 2. Download the dataset as a .zip file. 3. Extract the .zip file to a local directory.

The dataset contains multiple CSV files: - `studentInfo.csv`: Student demographics and performance data. - `studentVle.csv`: Interactions with the VLE. - `vle.csv`: Details of learning resources. - `studentAssessment.csv`: Assessment scores.

### Loading OULAD Data in R

Once the data is downloaded and extracted, follow these steps to load and access it in R:

#### Step 1: Install Required Packages

```
# Install necessary packages
install.packages(c("readr", "dplyr"))
```

## Step 2: Load Data

Use the `readr` package to read the CSV files into R.

```
# Load required libraries
library(readr)

# Define the path to the OULAD data
data_path <- "path/to/OULAD/"

# Load individual CSV files
student_info <- read_csv(file.path(data_path, "studentInfo.csv"))
student_vle <- read_csv(file.path(data_path, "studentVle.csv"))
vle <- read_csv(file.path(data_path, "vle.csv"))
student_assessment <- read_csv(file.path(data_path, "studentAssessment.csv"))
```

## Step 3: Preview the Data

Inspect the structure and contents of the datasets.

```
# View the first few rows of student info
head(student_info)

# Check the structure of the student VLE data
str(student_vle)
```

### 4.2.2 Learning Analytics

#### What is Learning Analytics?

**Learning Analytics (LA)** refers to the measurement, collection, analysis, and reporting of data about learners and their contexts. The primary goal of LA is to understand and improve learning processes by identifying patterns, predicting outcomes, and providing actionable insights to educators, institutions, and learners.

Key features of LA include:

- **Data Collection:** Gathering information from digital platforms such as learning management systems (LMS) or external assessments.
- **Analysis:** Using machine learning, statistical methods, or visualization tools to reveal trends and patterns.
- **Applications:** Supporting personalized learning, enhancing institutional decision-making, and improving curriculum design.

## **Applications of Learning Analytics in Big Data**

Learning analytics can be applied to large-scale educational datasets like **PISA**, **IPEDS**, and **OULAD** to uncover trends, predict outcomes, and guide interventions.

### **1. PISA Data and Learning Analytics**

- **What it offers:** Insights into international student performance in reading, math, and science, combined with contextual variables (e.g., socio-economic status).
- **LA Applications:**
  - Identifying key factors influencing performance across countries.
  - Predicting the impact of ICT use on student achievement.
  - Segmenting students into performance clusters for targeted interventions.

### **2. IPEDS Data and Learning Analytics**

- **What it offers:** U.S. institutional-level data on enrollment, graduation rates, tuition, and financial aid.
- **LA Applications:**
  - Analyzing trends in student demographics across institutions.
  - Predicting enrollment patterns based on historical data.
  - Benchmarking institutions to inform policymaking and funding decisions.

### **3. OULAD and Learning Analytics**

- **What it offers:** Rich data on student engagement with virtual learning environments (VLE), assessment scores, and demographic information.
- **LA Applications:**
  - Tracking student interactions with learning resources to predict course completion.
  - Modeling the relationship between VLE usage and final grades.
  - Detecting early warning signs for at-risk students based on engagement metrics.

## **Why Learning Analytics Matters**

The integration of **Learning Analytics** with big data enables researchers and practitioners to:

- **Personalize Learning:** Tailor educational experiences to meet individual needs.
- **Improve Retention:** Identify at-risk learners and implement timely interventions.
- **Enhance Decision-Making:** Provide evidence-based recommendations for curriculum and policy adjustments.

By leveraging datasets like PISA, IPEDS, and OULAD, learning analytics can help bridge the gap between raw data and actionable insights, fostering a more equitable and effective educational landscape.

## **Supervised Learning in Learning Analytics**

Machine Learning, particularly **Supervised Learning**, has become a cornerstone of Learning Analytics. Supervised learning models are trained on labeled datasets, where input features are mapped to known outcomes, enabling the prediction of new, unseen data.

### **Key Concepts in Supervised Learning**

- **Definition**

Supervised Learning is a subset of Machine Learning focused on learning a mapping between input variables (features) and output variables (labels or outcomes). Models trained on labeled data can predict outcomes for new data points.

- **Common Algorithms**

- Linear Regression
- Logistic Regression
- Decision Trees and Random Forests
- Neural Networks

- **Applications in Education**

Supervised learning is particularly effective in Learning Analytics for predicting:

- Student performance
- Dropout risks
- Enrollment trends
- Course completion rates

## **Applications of Supervised Learning with Big Data**

### **1. PISA Data and Supervised Learning**

- **Goal:** Use demographic and contextual features to predict student performance in mathematics, reading, or science.
- **Example:** Train a linear regression model to identify the relationship between socioeconomic status and test scores.

## **2. IPEDS Data and Supervised Learning**

- **Goal:** Develop models to predict institutional enrollment rates based on financial aid, demographics, and program offerings.
- **Example:** Use logistic regression to forecast whether a student is likely to enroll based on financial aid eligibility.

## **3. OULAD Data and Supervised Learning**

- **Goal:** Predict student outcomes (e.g., pass/fail) based on engagement metrics like forum participation and assignment submissions.
- **Example:** Train a random forest model to classify students as “at-risk” or “not at-risk” based on weekly interaction data.

### **Choosing the Right Supervised Learning Approach**

When applying supervised learning in Learning Analytics:

1. **Define the Goal:** Clearly articulate the outcome you want to predict (e.g., performance, enrollment, or engagement).
2. **Select an Algorithm:** Choose an appropriate model based on the data and prediction task.
  - For continuous outcomes, use regression models.
  - For categorical outcomes, use classification models like logistic regression or random forests.
3. **Feature Engineering:** Select and preprocess relevant features (e.g., attendance, demographics, assignment scores) to improve model accuracy.
4. **Evaluate Model Performance:** Use metrics such as accuracy, precision, recall, or R-squared to assess model effectiveness.

Integrating supervised learning techniques into Learning Analytics, researchers and practitioners can leverage big data to make data-driven predictions and decisions, ultimately enhancing educational outcomes.

## **4.3 Logistic Regression ML**

### **4.3.1 Purpose + CASE**

#### **Purpose**

Logistic regression is a supervised learning technique widely used for binary classification tasks. It models the probability of an event occurring (e.g., success vs. failure) based on a set of predictor variables. Logistic regression is particularly effective in educational research for predicting outcomes such as retention, enrollment, or graduation rates.

## CASE: Predicting Graduation Rates

This case study is based on IPEDS data and inspired by Zong and Davis (2022). We predict graduation rates as a binary outcome (`good_grad_rate`) using institutional features such as total enrollment, admission rate, tuition fees, and average instructional staff salary.

### 4.3.2 Sample Research Questions (RQs)

- **RQ A:** What institutional factors are associated with high graduation rates in U.S. four-year universities?
- **RQ B:** How accurately can we predict high graduation rates using institutional features with supervised machine learning?

### 4.3.3 Analysis

#### Loading Required Packages

We load necessary R packages for data wrangling, cleaning, and modeling.

```
# Load necessary libraries for data cleaning, wrangling, and modeling
library(tidyverse) # For data manipulation and visualization
library(tidymodels) # For machine learning workflows
library(janitor)    # For cleaning variable names
```

#### Loading and Cleaning Data

We read the IPEDS dataset and clean column names for easier handling.

```
# Read in IPEDS data from CSV file
ipeds <- read_csv("data/ipeds-all-title-9-2022-data.csv")

# Clean column names for consistency and usability
ipeds <- janitor::clean_names(ipeds)
```

#### Data Wrangling

Select relevant variables, filter the dataset, and create the dependent variable `good_grad_rate`.

```

# Select and rename key variables; filter relevant institutions
ipeds <- ipeds %>%
  select(
    name = institution_name, # Institution name
    total_enroll = drvref2022_total_enrollment, # Total enrollment
    pct_admitted = drvadm2022_percent_admitted_total, # Admission percentage
    tuition_fees = drvic2022_tuition_and_fees_2021_22, # Tuition fees
    grad_rate = drvgr2022_graduation_rate_total_cohort, # Graduation rate
    percent_fin_aid = sfa2122_percent_of_full_time_first_time_undergraduates_awarded_any_financial_aid,
    avg_salary =drvhr2022_average_salary_equated_to_9_months_of_full_time_instructional_staff_salary
  ) %>%
  filter(!is.na(grad_rate)) %>% # Remove rows with missing graduation rates
  mutate(
    # Create binary dependent variable for high graduation rates
    good_grad_rate = if_else(grad_rate > 62, 1, 0),
    good_grad_rate = as.factor(good_grad_rate) # Convert to factor
  )

```

## Exploratory Data Analysis (EDA)


Visualize the distribution of the graduation rate.

```

# Plot a histogram of graduation rates
ipeds %>%
  ggplot(aes(x = grad_rate)) +
  geom_histogram(bins = 20, fill = "blue", color = "white") +
  labs(
    title = "Distribution of Graduation Rates",
    x = "Graduation Rate",
    y = "Frequency"
  )

```

## Distribution of Graduation Rates



## Logistic Regression Model

Fit a logistic regression model to predict high graduation rates.

```
# Fit logistic regression model
m1 <- glm(
  good_grad_rate ~ total_enroll + pct_admitted + tuition_fees + percent_fin_aid + avg_salary,
  data = ipeds,
  family = "binomial" # Specify logistic regression for binary outcome
)

# View model summary
summary(m1)
```

Call:

```
glm(formula = good_grad_rate ~ total_enroll + pct_admitted +
    tuition_fees + percent_fin_aid + avg_salary, family = "binomial",
    data = ipeds)
```

Coefficients:

	Estimate	Std. Error	z value	Pr(> z )
(Intercept)	-8.742e-01	6.237e-01	-1.402	0.161

```

total_enroll      3.350e-05  7.880e-06   4.251 2.13e-05 ***
pct_admitted     -1.407e-02  3.519e-03  -3.997 6.40e-05 ***
tuition_fees     6.952e-05  4.965e-06   14.003 < 2e-16 ***
percent_fin_aid -2.960e-02  5.652e-03  -5.237 1.64e-07 ***
avg_salary       2.996e-05  3.870e-06   7.740 9.91e-15 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 2277  on 1706  degrees of freedom
Residual deviance: 1632  on 1701  degrees of freedom
(3621 observations deleted due to missingness)
AIC: 1644

Number of Fisher Scoring iterations: 5

```

## Supervised ML Workflow

Use the `tidymodels` framework to build a machine learning model.

```

# Define recipe for the model (preprocessing steps)
my_rec <- recipe(good_grad_rate ~ total_enroll + pct_admitted + tuition_fees + percent_fin_aid)

# Specify logistic regression model with tidymodels
my_mod <- logistic_reg() %>%
  set_engine("glm") %>%          # Use glm engine for logistic regression
  set_mode("classification")      # Specify binary classification task

# Create workflow to connect recipe and model
my_wf <- workflow() %>%
  add_recipe(my_rec) %>%
  add_model(my_mod)

# Fit the logistic regression model
fit_model <- fit(my_wf, ipeds)

# Generate predictions on the dataset
predictions <- predict(fit_model, ipeds) %>%
  bind_cols(ipeds) # Combine predictions with original data

# Calculate and display accuracy

```

```

my_accuracy <- predictions %>%
  metrics(truth = good_grad_rate, estimate = .pred_class) %>%
  filter(.metric == "accuracy")

my_accuracy

# A tibble: 1 x 3
  .metric   .estimator .estimate
  <chr>     <chr>        <dbl>
1 accuracy  binary      0.800

```

#### 4.3.4 Results and Discussions

##### Logistic Regression Model (RQ A)

The logistic regression model was fitted to predict whether a university achieves a “good” graduation rate (i.e., graduation rate  $> 62\%$ ) based on several institutional features. The model output is summarized below:

- **Coefficients & Significance:**
  - **total\_enroll:** Estimate = 3.35e-05,  $z = 4.251$ ,  $p = 2.13e-05$   
*Interpretation:* As total enrollment increases, the probability of a high graduation rate increases.
  - **pct\_admitted:** Estimate = -1.407e-02,  $z = -3.997$ ,  $p = 6.40e-05$   
*Interpretation:* Higher admission percentages are associated with a lower likelihood of achieving a high graduation rate.
  - **tuition\_fees:** Estimate = 6.952e-05,  $z = 14.003$ ,  $p < 2e-16$   
*Interpretation:* Higher tuition fees are strongly associated with higher graduation rates.
  - **percent\_fin\_aid:** Estimate = -2.960e-02,  $z = -5.237$ ,  $p = 1.64e-07$   
*Interpretation:* A higher percentage of students receiving financial aid is associated with a lower probability of a good graduation rate.
  - **avg\_salary:** Estimate = 2.996e-05,  $z = 7.740$ ,  $p = 9.91e-15$   
*Interpretation:* Higher average salaries for instructional staff are positively associated with high graduation rates.
- **Model Fit Statistics:**
  - **Null Deviance:** 2277 (on 1706 degrees of freedom)
  - **Residual Deviance:** 1632 (on 1701 degrees of freedom)
  - **AIC:** 1644
  - Note: 3621 observations were deleted due to missing values.

Overall, the regression model demonstrates that several institutional factors are statistically significant predictors of graduation rates. In particular, tuition fees and avg\_salary have a strong positive effect, while pct\_admitted and percent\_fin\_aid show negative associations.

### **Supervised ML Workflow Results (RQ B)**

Using the tidymodels framework, we built a logistic regression model as part of a supervised machine learning workflow. The performance metric obtained is as follows:

- **Accuracy:** 80.02%

This indicates that the machine learning model correctly classified approximately 80% of the institutions as having either a good or not good graduation rate, based on the selected predictors.

## **Overall Discussion**

- **Similarities between Approaches:**

- Both the traditional logistic regression and the tidymodels workflow identified key predictors that influence graduation rates, such as total enrollment, admission percentage, tuition fees, financial aid percentage, and average staff salary.
- Each approach provides valuable insights: the regression model offers detailed coefficient estimates and significance levels, while the tidymodels workflow emphasizes predictive accuracy.

- **Differences between Approaches:**

- **Interpretability vs. Predictive Performance:** The logistic regression output delivers interpretability through its coefficients and p-values, allowing us to understand the direction and magnitude of the relationships. In contrast, the supervised ML workflow focuses on achieving a robust predictive performance, evidenced by an 80% accuracy.
- **Handling of Data:** The traditional regression model summarizes the relationship between variables, whereas the ML workflow integrates data pre-processing, modeling, and validation into a cohesive framework.

In summary, our analyses indicate that institutional factors, particularly tuition fees and staff salaries, play a significant role in predicting graduation outcomes. The supervised ML approach, with an accuracy of around 80%, confirms the model's practical utility in classifying institutions based on graduation performance. Both methods complement each other, providing a comprehensive understanding of the underlying dynamics that drive graduation rates in higher education.

## 4.4 Random Forests ML on Interactions Data

In this section, we explore a more sophisticated supervised learning approach—Random Forests—to model student interactions data from the Open University Learning Analytics Dataset (OULAD). Building on our earlier work with logistic regression and evaluation metrics, this case study examines whether a random forest model can improve predictive performance when leveraging clickstream data from the virtual learning environment (VLE).

### 4.4.1 Purpose + CASE

#### Purpose

Random Forests is an ensemble learning method that builds multiple decision trees and aggregates their results to improve prediction accuracy and control over-fitting. It is particularly well suited for complex, high-dimensional data such as student interaction (clickstream) data. This approach not only provides robust predictions but also offers insights into variable importance, helping us understand which features most influence student outcomes.

#### CASE

Inspired by research on digital trace data (e.g., Rodriguez et al., 2021; Bosch, 2021), this case study uses pre-processed interactions data from OULAD. In our analysis, we focus on predicting whether a student will pass the course (a binary outcome) based on engineered features derived from clickstream data. These features include the total number of clicks (`sum_clicks`), summary statistics (mean and standard deviation of clicks), and linear trends over time (slope and intercept from clickstream patterns).

### 4.4.2 Sample Research Questions

- **RQ1:** How accurately can a random forest model predict whether a student will pass a course using interactions data from OULAD?
- **RQ2:** Which interaction-based features (e.g., total clicks, click stream slope) are most important in predicting student outcomes?
- **RQ3:** How does the use of cross-validation (e.g., v-fold CV) influence the stability and generalizability of the random forest model on interactions data?

#### 4.4.3 Analysis

##### Loading Required Packages

```
# Load necessary libraries for data manipulation and modeling
library(tidyverse)      # Data wrangling and visualization
library(janitor)        # Cleaning variable names
library(tidymodels)     # Modeling workflow
library(ranger)         # Random forest implementation
library(vip)             # Variable importance plots
```

##### Loading and Preparing the Data

We load the pre-filtered interactions data from OULAD along with a students-and-assessments file, then join them to create a complete dataset for modeling.

```
# Load the interactions data (filtered for the first one-third of the semester)
interactions <- read_csv("data/oulad-interactions-filtered.csv")

# Load the students and assessments data
students_and_assessments <- read_csv("data/oulad-students-and-assessments.csv")

# Create cut-off dates based on assessments data (using first quantile as intervention point)
assessments <- read_csv("data/oulad-assessments.csv")

# Create cut-off dates based on assessments data using the correct date column 'date_submitted'
code_module_dates <- assessments %>%
  group_by(code_module, code_presentation) %>%
  summarize(quintile_cutoff_date = quantile(date_submitted, probs = 0.25, na.rm = TRUE), .by_group = TRUE)

# Join interactions with the cutoff dates and filter
interactions_joined <- interactions %>%
  left_join(code_module_dates, by = c("code_module", "code_presentation"))

interactions_joined <- interactions_joined %>%
  select(-quintile_cutoff_date.x) %>%
  rename(quintile_cutoff_date = quintile_cutoff_date.y)
```

```

# Filter interactions to include only those before the cutoff date
interactions_filtered <- interactions_joined %>%
  filter(date < quantile_cutoff_date)

# Summarize interactions: total clicks, mean and standard deviation
interactions_summarized <- interactions_filtered %>%
  group_by(id_student, code_module, code_presentation) %>%
  summarize(
    sum_clicks = sum(sum_click),
    sd_clicks = sd(sum_click),
    mean_clicks = mean(sum_click)
  )

# (Optional) Further feature engineering: derive linear slopes from clickstream
fit_model <- function(data) {
  tryCatch(
    {
      model <- lm(sum_click ~ date, data = data)
      tidy(model)
    },
    error = function(e) { tibble(term = NA, estimate = NA, std.error = NA, statistic = NA),
    warning = function(w) { tibble(term = NA, estimate = NA, std.error = NA, statistic = NA)
    }
  }
}

interactions_slopes <- interactions_filtered %>%
  group_by(id_student, code_module, code_presentation) %>%
  nest() %>%
  mutate(model = map(data, fit_model)) %>%
  unnest(model) %>%
  ungroup() %>%
  select(code_module, code_presentation, id_student, term, estimate) %>%
  filter(!is.na(term)) %>%
  pivot_wider(names_from = term, values_from = estimate) %>%
  mutate_if(is.numeric, round, 4) %>%
  rename(intercept = `Intercept`, slope = date)

# Join summarized clicks and slopes features
interactions_features <- left_join(interactions_summarized, interactions_slopes, by = c("id_s",
  "code_module", "code_presentation"))

# Finally, join with students_and_assessments to get the outcome variable
students_assessments_and_interactions <- left_join(students_and_assessments, interactions_fea

```

```

# Ensure outcome variable 'pass' is a factor
students_assessments_and_interactions <- students_assessments_and_interactions %>%
  mutate(pass = as.factor(pass))

# Optional: Inspect the final dataset
students_assessments_and_interactions %>%
  skimr::skim()

```

Table 3: Data summary

Name	Piped data
Number of rows	32593
Number of columns	22
<hr/>	
Column type frequency:	
character	8
factor	1
numeric	13
<hr/>	
Group variables	None

#### Variable type: character

skim_variable	n_missing	complete_rate	min	max	empty	n_unique	whitespace
code_module	0	1	3	3	0	7	0
code_presentation	0	1	5	5	0	4	0
gender	0	1	1	1	0	2	0
region	0	1	5	20	0	13	0
highest_education	0	1	15	27	0	5	0
age_band	0	1	4	5	0	3	0
disability	0	1	1	1	0	2	0
final_result	0	1	4	11	0	4	0

#### Variable type: factor

skim_variable	n_missing	complete_rate	ordered	n_unique	top_counts
pass	0	1	FALSE	2	0: 20232, 1: 12361

#### Variable type: numeric

skim_variable	n_missing	complete_rate	sd	p0	p25	p50	p75	p100	hist
id_student	0	1.00	706687.6749167.3733.0008573.090310.0014453.0716795.00						
imd_band	4627	0.86	5.62	2.73	1.00	4.00	6.00	8.00	10.00
num_of_prev_attempts	1.00	0.16	0.48	0.00	0.00	0.00	0.00	0.00	6.00
studied_credits	0	1.00	79.76	41.07	30.00	60.00	60.00	120.00	655.00
module_presentation_length	1.00	256.01	13.18	234.00	241.00	262.00	268.00	269.00	
date_registration	45	1.00	-	49.26	-	-	-	-	167.00
				69.41		322.00	100.00	57.00	29.00
date_unregistration	29521	0.31	49.76	82.46	-	-2.00	27.00	109.00	444.00
						365.00			
mean_weighted_score	7058	0.76	544.70	381.39	0.00	160.00	610.00	875.00	1512.00
sum_clicks	3495	0.89	474.93	572.89	1.00	128.00	295.50	604.00	10712.00
sd_clicks	3753	0.88	4.91	5.51	0.00	2.37	3.72	6.44	560.24
mean_clicks	3495	0.89	3.19	1.30	1.00	2.33	2.95	3.82	47.12
intercept	3640	0.89	3.04	4.61	-	2.15	2.80	3.66	130.83
						585.59			
slope	4441	0.86	0.01	0.22	-	-0.01	0.01	0.03	20.12
						12.17			

## Creating the Model Recipe

We build a recipe that includes the engineered features from interactions data along with other predictors from the students data.

```
my_rec2 <- recipe(pass ~ disability +
  date_registration +
  gender +
  code_module +
  mean_weighted_score +
  sum_clicks + sd_clicks + mean_clicks +
  intercept + slope,
  data = students_assessments_and_interactions) %>%
  step_dummy(disability) %>%
  step_dummy(gender) %>%
  step_dummy(code_module) %>%
  step_impute_knn(mean_weighted_score) %>%
  step_impute_knn(sum_clicks) %>%
  step_impute_knn(sd_clicks) %>%
  step_impute_knn(mean_clicks) %>%
  step_impute_knn(intercept) %>%
  step_impute_knn(slope) %>%
```

```
step_impute_knn(date_registration) %>%
  step_normalize(all_numeric_predictors())
```

## Specifying the Model and Workflow

We use a random forest model via the `ranger` engine and set up our workflow.

```
# Specify random forest model
my_mod2 <- rand_forest() %>%
  set_engine("ranger", importance = "impurity") %>%
  set_mode("classification")

# Create workflow to bundle the recipe and model
my_wf2 <- workflow() %>%
  add_recipe(my_rec2) %>%
  add_model(my_mod2)
```

## Resampling and Model Fitting

We perform cross-validation (v-fold CV) to estimate model performance.

```
# Create 4-fold cross-validation on training data
vfcv <- vfold_cv(data = students_assessments_and_interactions, v = 4, strata = pass)

# Specify metrics: accuracy, sensitivity, specificity, ppv, npv, and Cohen's kappa
class_metrics <- metric_set(accuracy, sensitivity, specificity, ppv, npv, kap)

# Fit the model using resampling
fitted_model_resamples <- fit_resamples(my_wf2, resamples = vfcv, metrics = class_metrics)

# Collect and display metrics
collect_metrics(fitted_model_resamples)
```

.metric	.estimator	mean	n	std_err	.config
<chr>	<chr>	<dbl>	<int>	<dbl>	<chr>
1 accuracy	binary	0.671	4	0.000795	Preprocessor1_Model1
2 kap	binary	0.263	4	0.00351	Preprocessor1_Model1
3 npv	binary	0.589	4	0.000393	Preprocessor1_Model1
4 ppv	binary	0.703	4	0.00175	Preprocessor1_Model1

```

5 sensitivity binary      0.814      4 0.00293 Preprocessor1_Model1
6 specificity binary     0.436      4 0.00680 Preprocessor1_Model1

```

## Final Model Fit and Evaluation

Finally, we fit the model on the full training set (using `last_fit`) and evaluate its predictions on the test set.

```

# Split data into training and testing sets (e.g., 33% for testing)
set.seed(20230712)
train_test_split <- initial_split(students_assessments_and_interactions, prop = 0.67, strata = "student_id")
data_train <- training(train_test_split)
data_test <- testing(train_test_split)

# Fit final model on the training set and evaluate on the test set
final_fit <- last_fit(my_wf2, train_test_split, metrics = class_metrics)

# Collect and display final metrics
collect_metrics(final_fit)

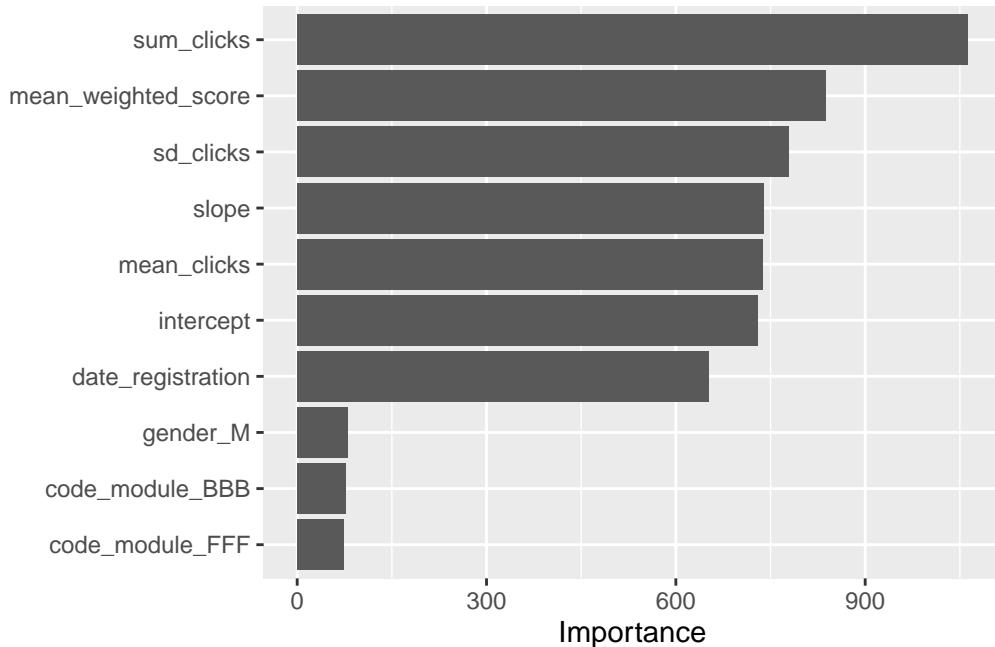
# A tibble: 6 x 4
#> #> .metric      .estimator .estimate .config
#> #> <chr>        <chr>       <dbl> <chr>
#> 1 accuracy    binary      0.665 Preprocessor1_Model1
#> 2 sensitivity binary      0.815 Preprocessor1_Model1
#> 3 specificity binary      0.419 Preprocessor1_Model1
#> 4 ppv         binary      0.697 Preprocessor1_Model1
#> 5 npv         binary      0.580 Preprocessor1_Model1
#> 6 kap         binary      0.247 Preprocessor1_Model1

# Generate and display a confusion matrix for final predictions
collect_predictions(final_fit) %>%
  conf_mat(.pred_class, pass)

  Truth
Prediction   0     1
          0 5439 1238
          1 2370 1710

```

```
# Extract the fitted model from the final workflow and plot variable importance
final_fit %>%
  pluck(".workflow", 1) %>% # Extract the workflow object from the final fit
  extract_fit_parsnip() %>% # Retrieve the fitted model from the workflow
  vip(num_features = 10)      # Plot the top 10 important features
```



```
# Extract the fitted model from the workflow
final_model <- final_fit %>%
  pluck(".workflow", 1) %>%
  extract_fit_parsnip()
# Extract the variable importance values from the fitted model
importance_values <- final_model$fit$variable.importance

# Print the variable importance values
print(importance_values)
```

date_registration	mean_weighted_score	sum_clicks	sd_clicks
652.14806	838.06052	1062.49141	778.63848
mean_clicks	intercept	slope	disability_Y
737.79417	729.52344	739.85682	60.93367
gender_M	code_module_BBB	code_module_CCC	code_module_DDD
79.62440	77.18477	70.72006	46.68243

code_module_EEE	code_module_FFF	code_module_GGG
28.26595	74.31310	35.27026

#### 4.4.4 Results and Discussions

##### Research Question 1 (RQ1):

*How accurately can a random forest model predict whether a student will pass a course using interactions data from OULAD?*

##### Response:

Using 4-fold cross-validation, our random forest model yielded an average accuracy of approximately **67.0%** (mean accuracy from resamples: 0.670) with a Cohen's Kappa of **0.261**, suggesting moderate agreement beyond chance. When fitted on the entire training set and evaluated on the test set, the final model showed an accuracy of **66.5%** along with: - **Sensitivity:** 81.5% – indicating the model correctly identifies a high proportion of students who pass. - **Specificity:** 41.9% – suggesting that the model is less effective at correctly identifying students who do not pass. - **Positive Predictive Value (PPV):** 69.7% - **Negative Predictive Value (NPV):** 58.0%

The confusion matrix shows: - True Negatives (TN): 5439 - False Negatives (FN): 1238 - False Positives (FP): 2370 - True Positives (TP): 1710

Overall, these metrics indicate that while the model performs well in detecting positive outcomes (high sensitivity), its lower specificity means that it tends to misclassify a relatively higher proportion of non-passing students.

##### Research Question 2 (RQ2):

*Which interaction-based features are most important in predicting student outcomes?*

##### Response:

The variable importance analysis, extracted from the final random forest model using the `vip()` function, highlights the following key predictors (with their respective importance scores):

- **sum\_clicks:** 1062.49 – This is the most influential feature, indicating that the total number of clicks (i.e., student engagement) in the VLE is a strong predictor of student success.
- **mean\_weighted\_score:** 838.06 – Reflecting academic performance as measured by weighted assessment scores.
- **mean\_clicks:** 737.79, **slope:** 739.86, and **intercept:** 729.52 – These engineered features representing the central tendency and trend of click behavior further underline the importance of digital engagement patterns.

- **date\_registration:** 652.15 – The registration date also plays a significant role.
- Other categorical variables (e.g., dummy-coded **disability**, **gender**, and **code\_module** levels) generally show lower importance scores, with values typically under 80, indicating that while they do contribute, engagement and performance metrics dominate.

These results suggest that both the intensity and the temporal trend of student interactions with the learning environment are critical in predicting whether a student will pass.

### **Research Question 3 (RQ3):**

*How does the use of cross-validation impact the stability and generalizability of the random forest model on interactions data?*

#### **Response:**

The use of 4-fold cross-validation (via `vfold_cv`) allowed us to assess the model's performance across multiple subsets of the data, mitigating the risk of overfitting. The resampling results are relatively consistent (with accuracy around 67%, sensitivity at 81.6%, and specificity around 43.2%), which supports the model's robustness and generalizability. Although the final test set performance (accuracy of 66.5%) is slightly lower, the overall consistency of metrics across folds indicates that our model is stable when applied to unseen data.

## **Overall Discussion**

The random forest model built on interactions data from OULAD demonstrates decent predictive performance with an accuracy of approximately 66.5–67% and high sensitivity (around 81.5%), indicating strong capability in identifying students who will pass the course. However, the relatively low specificity (around 42%) suggests that there is room for improvement in correctly classifying students who are at risk of not passing.

The variable importance analysis underscores that engagement-related features—especially **sum\_clicks** and features capturing the trend in interactions (slope, **mean\_clicks**)—are the most influential predictors. This insight implies that the digital footprint of student engagement in the virtual learning environment is critical for predicting academic outcomes.

In summary, while our model performs robustly across cross-validation folds and provides actionable insights into key predictive features, the lower specificity points to the need for further refinement. Future work might explore additional feature engineering, alternative model tuning, or combining models to better balance sensitivity and specificity, ultimately supporting timely interventions in educational settings.

# **Section 5 Multimodal Data (Images, Video, Audio) with LLMs**

## **5.1 Overview**

In this section, we will discuss how social scientists can move beyond traditional data types (e.g., text and numbers) and learn about capturing and analyzing multimodal data.

Multimodal data includes audio, video, and other non-textual information that gives a fuller picture of human behavior.

Modern devices such as wearables, smartphones, and online platforms now let researchers collect large amounts of this mixed data.

To make sense of it, we use computational tools that combine the different types:

- image analysis for video frames,
- voice-to-text software for audio, and
- machine-learning models that link text, pictures, and sensor signals.

These tools let researchers ask new questions. For example, how body language and tone of voice together affect conversations or how physiological signals match feelings can help uncover insights that single-mode studies miss. By integrating multimodal data, social scientists can broaden the depth and reach of their research beyond what conventional single-mode analysis offers.

## **5.2 Images**

Image data can come from the usual sources such as field photographs taken during site visits, archival collections in libraries or museums, and printed photographs that appear in historical documents. Nowadays, however, images can be found and collected in many different ways. For example, social media platforms like Instagram, Facebook, and TikTok are rich with

user-generated photos; online photo repositories such as Flickr, Unsplash, and Wikimedia Commons host millions of images that are freely accessible; news outlets regularly publish photographs to accompany stories; satellite imagery from NASA or ESA provides large-scale visual data; and everyday smartphone cameras capture images that can be shared in research settings. Please note that although we cover some prominent ways, this is by no means an exhaustive list. Therefore, please refer to the additional resources section at the end of the section to dive deeper.

### 3.3 Analyzing Images

With the advent of Large-Language Models (LLMs) we can use their power to analyze images. In this section, we will focus on using one package that uses local LLMs (i.e., privacy) to analyze image files: [{kuzco}](#).

Kuzco is

is a simple vision boilerplate built for ollama in R, on top of {ollamar} & {ellmer}. {kuzco} is designed as a computer vision assistant, giving local models guidance on classifying images and return structured data. The goal is to standardize outputs for image classification and use LLMs as an alternative option to keras or torch. {kuzco} currently supports: classification, recognition, sentiment, text extraction, alt-text creation, and **custom** computer vision tasks.

#### 5.3.1. Setting Up Kuzco

To use kuzco, you need to, first, install Ollama (a software that allows pulling and running local LLMs) and ollamar & ellmer packages.

You can install Ollama by downloading and installing the application from its provider's website. Basically the steps are:

1. Download and install the [Ollama](#) app.
  - [macOS](#)
  - [Windows preview](#)
  - Linux: `curl -fsSL https://ollama.com/install.sh | sh`
  - [Docker image](#)
2. Open/launch the Ollama app to start the local server.

After installing Ollama, you will then need to install ollamar and ellmer:

```
install.packages("ollamar")
install.packages("ellmer")
```

Once these are installed, install kuzco:

```
devtools::install_github("frankiethull/kuzco")
```

### 5.3.2 Image Classification

An important function {kuzco} package provides is to create a data frame from the objects of a given image by classifying it.

#### Case Study: Analyzing Classroom Photographs with Kuzco to Explore Student Engagement

##### 5.3.2.1 Purpose

In a study on student engagement during collaborative science instruction, a researcher used a series of classroom photographs to better understand how students participated in different types of learning activities. Rather than relying solely on manual observation and field notes, the researcher applied the {kuzco} R package to process and interpret visual data. Three key functions—`l1m_image_classification()`, `l1m_image_sentiment()`, and `l1m_image_recognition()`—were used to generate insights about classroom scenes.

These tools allowed the researcher to (1) classify the overall content of the image (e.g., lab work, discussion, presentation), (2) recognize and count key objects or people in the frame (e.g., students, materials, whiteboards), and (3) estimate the emotional tone of the scene based on posture and facial cues. This approach enabled a more systematic and scalable analysis of classroom engagement, providing structured outputs that could be interpreted alongside observational data and interview responses.

##### 5.3.2.2 Research Questions

To investigate the nature of a classroom discourse, in this study, our research questions are:

- **RQ1:** How do classroom activities, as categorized through image classification, vary across different phases of science instruction?

- **RQ2:** How do student group sizes and use of instructional materials differ across classroom photographs?
- **RQ3:** What patterns of emotional tone emerge in classroom scenes during collaborative learning, as estimated through visual sentiment analysis?

### 5.3.2.3 Methods

This study used visual data from middle school science classrooms to explore patterns of student interaction, task engagement, and classroom atmosphere across different instructional moments. The analysis was supported by large language model (LLM)-based image processing tools from the `{kuzco}` R package, allowing for efficient classification, recognition, and sentiment estimation without advanced machine learning expertise.

### 5.3.2.4 Data Source

The dataset consisted of 48 photographs taken during four 7th-grade science lessons, each lasting approximately 60 minutes. Photos were captured every 5–7 minutes by a stationary camera positioned at the back of the room to minimize disruption. All images were de-identified prior to analysis to protect student privacy. Each photo represented a naturally occurring moment of group-based learning and was accompanied by a brief instructional context log maintained by the classroom observer.

### 5.3.2.5 Data Analysis

Images were processed using the following `{kuzco}` functions:

- `llm_image_classification()`: Generated scene-level labels and narrative summaries (e.g., “students engaged in group discussion around lab materials”).
- `llm_image_recognition()`: Identified and counted key visual entities such as students, desks, instructional materials, and gestures
- `llm_image_sentiment()`: Estimated the emotional tone of each scene (e.g., positive, neutral, frustrated), with particular attention to student posture and interaction dynamics.

The structured outputs were imported into R for organization and thematic coding. Using both deductive categories (e.g., group size, task type) and inductive patterns (e.g., collaborative vs. passive positioning), the researcher examined how engagement varied across activities. Triangulation with field notes enhanced interpretive validity, and descriptive summaries were generated to visualize classroom dynamics over time.

*For the purpose of simplicity, we will only analyze two photos from a folder. The process for batch analysis can be increased to more photos.*



With the code below, we create a function to batch analyze images:

```
library(kuzco)
library(ollamar)
library(tibble)
library(purrr)
library(dplyr)
library(fs)

# Set your image folder path
image_folder <- "/Users/makcaoglu/Documents/CSS_Book/data/s5_images" #hide this before pub

# List images (adjust pattern as needed)
image_files <- dir_ls(image_folder, regexp = "\\\.(jpg|jpeg|png)$", recurse = FALSE)

# Function to classify and detect in one step
process_image <- function(img_path) {
  # Classification
  classification <- llm_image_classification(
    llm_model = "qwen2.5vl",
    image = img_path,
    backend = 'ollamar'
  )

  # Object detection (e.g., people)
  detection <- llm_image_recognition(
    llm_model = "qwen2.5vl",
    image = img_path,
    recognize_object = "people",
    backend = 'ollamar'
  )

  # Sentiment/emotion
  sentiment <- llm_image_sentiment(
    llm_model = "qwen2.5vl",
    image = img_path
  )

  #the new custom fuction for sentiment
  customized <- llm_image_custom(
    llm_model = "qwen2.5vl",
    image = img_path,
    system_prompt = "You are an expert classroom observer. You analyze classroom photographs"
}
```

```

image_prompt = "Describe the overall sentiment of the classroom and explain what visual elements support this sentiment"
example_df = data.frame(
  classroom_sentiment = "positive",
  engagement_level = "high",
  sentiment_rationale = "Students are smiling, interacting with each other, and appear attentive"
)
)

# Return combined tibble
tibble::tibble(
  file = img_path,
  image_classification = classification$image_classification,
  primary_object = classification$primary_object,
  secondary_object = classification$secondary_object,
  image_description = classification$image_description,
  image_colors = classification$image_colors,
  image_proba_names = paste(unlist(classification$image_proba_names), collapse = ", "),
  image_proba_values = paste(unlist(classification$image_proba_values), collapse = ", "),
  object_recognized = detection$object_recognized,
  object_count = detection$object_count,
  object_description = detection$object_description,
  object_location = detection$object_location,
  classroom_sentiment = customized$classroom_sentiment,
  engagement_level = customized$engagement_level,
  sentiment_rationale = customized$sentiment_rationale
)
}

}

```

Now, we run the analyses:

```

# Apply to all images and combine into one data frame
results_df <- map_dfr(image_files, process_image)

# View result
print(results_df)

# Arrange columns in logical order and rename for clarity
results_clean <- results_df |>
  select(
    file,

```

```

    image_classification,
    image_description,
    primary_object,
    secondary_object,
    image_colors,
    image_proba_names,
    image_proba_values,
    object_recognized,
    object_count,
    object_description,
    object_location,
    image_sentiment,
    image_score,
    sentiment_description,
    image_keywords
  )

# Save to CSV (optional)
write.csv(results_clean, "image_classification_detection_results.csv", row.names = FALSE)

# View top images with the most people (if desired)
results_clean |>
  arrange(desc(object_count)) |>
  print(n = Inf)

```

### 5.3.2.6 Results and Discussion:

The analysis of classroom photographs using the `{kuzco}` package yielded structured insights across three domains: instructional context (classification), observable features (recognition), and affective tone (sentiment). Below, we summarize preliminary findings from two sample images.

#### RQ1: Instructional Activity Patterns (Image Classification)

Across images, `l1m_image_classification()` consistently labeled the scenes as “classroom,” with descriptive summaries indicating quiet, focused individual or small-group work. For example, one image was described as “a classroom scene with a student focused on writing,” while another featured “students sitting at desks... and a teacher in the background.” These classifications help distinguish between instructional moments (e.g., individual seatwork vs. whole-group interaction), informing a timeline of activity types over the course of the lesson.

#### RQ2: Student Presence and Interaction (Image Recognition)

The `l1m_image_recognition()` function identified and described key objects and individuals

in each photo. Object counts ranged from 2 to 7, allowing for comparisons in group size and density. In one image, a single student was the focal point, while another captured a larger group of students engaged in various literacy tasks. Object descriptions included not only counts but also contextual roles (e.g., “a boy writing with a pencil” or “students engaged in reading”), enhancing interpretability of engagement.

**RQ3: Affective Tone in Classroom Scenes (Image Sentiment)**

Both images were rated as **positive** using `llm_image_sentiment()`, with scores of 0.7 and 0.8 respectively. Sentiment descriptions highlighted emotional tone inferred from posture, interaction, and environmental cues: one described “a quiet studious atmosphere,” while the other conveyed “a cheerful and engaged sentiment.” These outputs suggest that students were generally attentive and involved, and that the physical learning environment contributed to a positive classroom climate.

# Section 6 Local LLMs in Educational Research

## **Abstract:**

The use of large language models (LLMs) in data analysis is rapidly increasing across education and social science research. However, concerns about data privacy, institutional data protection policies, and strict IRB (Institutional Review Board) procedures present significant challenges when using cloud-based or proprietary AI services. To address these challenges, this chapter introduces local LLM solutions—focusing on LM Studio—which allow researchers to run powerful models entirely on their own computers, ensuring data stays private and analysis remains flexible.

## **6.1 What are Local LLMs?**

Local LLMs are large language models that run directly on your own computer, rather than in the cloud. By processing data locally, they help ensure privacy, data sovereignty, and compliance with institutional or governmental regulations. Local LLMs can be open-source (such as Llama, Qwen, DeepSeek, Mistral) and are compatible with various operating systems and hardware.

**Key advantages of local LLMs:** - Data never leaves your computer - No need for external API keys or internet access to analyze sensitive data - Flexibility to use custom or open-source models - Often no usage fees

## **6.2 What Can Local LLMs Do?**

With the right setup, local LLMs can:

- Summarize, paraphrase, and analyze text data (open-ended survey responses, interview transcripts, etc.)
- Support qualitative and quantitative educational research workflows
- Generate coding frameworks, extract themes, or automate report writing
- Perform document-based question answering (“chat with your PDFs”)
- Integrate with other research tools via REST APIs

## 6.3 Getting Started with LM Studio

LM Studio is a free, cross-platform application that enables researchers to run, manage, and interact with local LLMs (such as Llama, DeepSeek, Qwen, Mistral, and gpt-oss) entirely on their own computers. By using LM Studio, you gain powerful, offline data analysis capabilities without sacrificing data privacy or compliance.

**Key Points:** - **Supported Platforms:** macOS (Apple Silicon), Windows (x64/ARM64), and Linux (x64). - **System Requirements:** For best results, consult the [System Requirements](#) page for recommended RAM, CPU/GPU, and storage.

### 6.3.1 Installation Steps

1. **Download LM Studio** for your operating system from the [official Downloads page](#).
2. **Install and launch** the application.
3. **Download your preferred LLM model** (such as Llama 3, Qwen, Mistral, DeepSeek, or gpt-oss) directly from within LM Studio.
4. **(Optional)** To use the API for scripting/automation, enable API access within LM Studio.
5. **(Optional)** Attach documents for “Chat with Documents” (RAG-style analysis) entirely offline.

#### Official Documentation:

- [LM Studio Docs - Getting Started Guide](#)

### 6.3.2 Main Features

- **Run local models** including Llama, Qwen, DeepSeek, Mistral, gpt-oss, and more.
- **Simple chat interface** for prompt-based interaction.
- **Offline “Chat with Documents”** for Retrieval Augmented Generation (RAG) use cases.
- **Search and download new models** from Hugging Face and other model hubs within LM Studio.
- **Manage models, prompts, and configurations** through a user-friendly GUI.
- **Serve local models** on OpenAI-compatible REST API endpoints, usable by R, Python, or other apps.
- **MCP server/client support** for advanced use cases.

### 6.3.3 API Integration

LM Studio exposes a REST API fully compatible with the OpenAI standard. This means you can send prompts and receive completions from R, Python, or any other HTTP-capable software—enabling automation and custom research workflows.

#### Example: Calling the LM Studio API from R

LM Studio exposes a REST API compatible with the OpenAI API standard. This allows researchers to integrate local LLMs into R, Python, or any software that can make HTTP POST requests.

```
library(httr) library(jsonlite)

prompt <- "Summarize the following open-ended survey responses: ..."

response <- POST( url = "http://localhost:1234/v1/completions",
                  body = toJSON(list( prompt = prompt,
                                      max_tokens = 200 ),
                  auto_unbox = TRUE),
                  encode = "json" )
content(response)
```

### 6.3.4 Summary Table of LM Studio Capabilities:

Feature	Description
Local LLMs	Run Llama, DeepSeek, Qwen, Mistral, etc. fully offline on your own machine
Chat Interface	Flexible prompt-based interaction
Document Chat (RAG)	Offline “chat with your documents”
Model Management	Download, organize, and switch between models
API Access	OpenAI-compatible REST endpoints for use with R, Python, scripts, apps
MCP Integration	Connect with and use MCP servers
Community & Support	<a href="#">Discord</a> , official docs, active development

## 6.4 Case Study: Comparing Local LLM Analysis to Traditional NLP on University AI Policy Texts

### 6.4.1 Research Question

Can a local LLM running via LM Studio reliably identify key themes in university AI policy statements—using the **same dataset** analyzed in Section 2—so that we can compare its results against traditional NLP methods and human coding?

---

### 6.4.2 Data Context

We reuse the **AI policy statements** dataset from Section 2, now simplified for privacy. The table has **one column** only:

- **Stance** (character): policy text (no institution names)

A typical structure (as seen in Section 2):

We will extract the same raw text field (**Stance**) so results are directly comparable to Section 2.

```
library(dplyr)
library(stringr)
library(readr)

# If 'university_policies' already exists (from Section 2), use it directly.
# Otherwise, safely fall back to reading the same CSV used in Section 2.
if (!exists("university_policies")) {
  university_policies <- read_csv("University_GenAI_Policy_Stance.csv", show_col_types = FALSE)
}

stopifnot("Stance" %in% names(university_policies))

policy_texts <- university_policies$Stance %>%
  as.character() %>%
  stringr::str_squish() %>%
  na.omit()

length(policy_texts)
```

```
[1] 99
```

```
head(policy_texts, 3)
```

```
[1] "If the text generated by ChatGPT is used as a starting point for original research or w  
[2] "Has ASU considered a ban on AI tools like other institutions such as NYU? No. ASU facul  
[3] "The following sample statements should be taken as starting points to craft your own po
```

#### 6.4.3 Implementation with LM Studio (Thematic Analysis)

We send the **same policy texts** to LM Studio's local API using the parameters already defined in your setup (`api_base`, `model_name`).

The model `openai/gpt-oss-20b` runs locally in LM Studio and provides OpenAI-compatible endpoints.

```
library(httr)
library(jsonlite)
library(glue)
library(stringr)

# Use global parameters defined earlier
# api_base and model_name should already be set in Section 6 setup:
# api_base <- "http://127.0.0.1:1234/v1"
# model_name <- "openai/gpt-oss-20b"

# ----- 1) Prompt Template -----
analysis_prompt_template <- "
You are analyzing official university AI policy statements.
Your task is to identify 3-5 key themes across the statements and report them in the exact fo

**INPUT DATA:**
- **Number of Statements:** {n_items}
- **Policy Statements:** {items}

**YOUR TASK:**
1) Identify 3-5 key themes across the policy statements.
2) For each theme:
   a) Provide a concise theme name.
   b) Provide a 1-2 sentence description.
   c) Provide one short verbatim example quote.
```

```

d) Provide an integer Frequency (count of statements mentioning it).
e) Provide Relative Frequency as a whole-number percentage.
3) Write a 3-5 sentence **Summary of Responses** synthesizing the most important insights.
4) Output strictly in the following format:

**Summary of Responses**
[3-5 sentence narrative summary goes here.]

**Thematic Table**
| Theme | Description | Illustrative Example(s) | Frequency | Relative Frequency |
|---|---|---|---|
| [Theme 1] | [Description] | - \"[Quote]\\" | [n] | [p]% |
| [Theme 2] | [Description] | - \"[Quote]\\" | [n] | [p]% |
"

# ----- 2) Chunk the corpus to stay within model context window -----
CHUNK_SIZE <- 15
chunks <- split(policy_texts, ceiling(seq_along(policy_texts) / CHUNK_SIZE))

# ----- 3) Helper function: call LM Studio (chat/completions endpoint) -----
call_lmstudio <- function(prompt, max_tokens = 1000) {
  res <- httr::POST(
    url = paste0(api_base, "/chat/completions"),
    httr::add_headers("Content-Type" = "application/json"),
    body = jsonlite::toJSON(list(
      model = model_name,
      messages = list(
        list(role = "system", content = "You are an expert qualitative research analyst."),
        list(role = "user", content = prompt)
      ),
      temperature = 0.2,
      max_tokens = max_tokens
    )), auto_unbox = TRUE)
  )
  httr::stop_for_status(res)
  content(res)$choices[[1]]$message$content
}

# ----- 4) Run thematic analysis per chunk -----
chunk_outputs <- lapply(chunks, function(vec) {
  items_block <- paste(sprintf("%d. %s", seq_along(vec), vec), collapse = "\n")
  final_prompt <- glue(analysis_prompt_template,

```

```

        n_items = length(vec),
        items   = items_block)
call_lmstudio(final_prompt)
})

# ----- 5) Merge all chunk-level analyses into a meta-synthesis -----
meta_prompt <- "
You will synthesize multiple chunk-level thematic analyses of the same corpus of university A
Unify and deduplicate themes across chunks, and output a single consolidated section in the order they appear in the input.

**Summary of Responses**
[3-5 sentence narrative summary.]

**Thematic Table**
| Theme | Description | Illustrative Example(s) | Frequency | Relative Frequency |
|---|---|---|---|
| [Unified Theme 1] | [Description] | - \"[Quote]\\" | [n] | [p]% |
| [Unified Theme 2] | [Description] | - \"[Quote]\\" | [n] | [p]% |
"
# Pairwise synthesis to reduce token usage
pairs <- split(chunk_outputs, ceiling(seq_along(chunk_outputs) / 2))

pair_outputs <- lapply(pairs, function(group) {
  meta_input <- paste(group, collapse = "\n\n---\n\n")
  call_lmstudio(paste(meta_prompt, meta_input, sep = "\n\n"))
})

# Now you have fewer intermediate syntheses
final_meta_input <- paste(pair_outputs, collapse = "\n\n---\n\n")
meta_output <- call_lmstudio(paste(meta_prompt, final_meta_input, sep = "\n\n"))
cat(meta_output)

#saveRDS(meta_output, "data/meta_output_saved.rds")
#saveRDS(meta_output, "data/meta_output_saved.rds")

```

```

library(stringr)
library(readr)

# --- Read RDS ---
meta_output <- readRDS("data/meta_output_saved.rds")

# --- Combine all elements into one long text block ---

```

```

meta_output_text <- paste(meta_output, collapse = "\n")

# --- Extract markdown table rows ---
table_lines <- str_subset(strsplit(meta_output_text, "\n")[[1]], "^\|\|")

# --- Clean leading/trailing pipes ---
table_text <- gsub("^\|\|", "", table_lines)

# --- Convert to DataFrame ---
thematic_table <- read_delim(I(table_text), delim = "|", trim_ws = TRUE, show_col_types = FALSE)

# --- Display result ---
print(thematic_table)

```

Theme	Description	Illustrative Example <sup>i</sup>	Frequency	Relative Frequency
<chr>	<chr>	<chr>	<chr>	<chr>
1 ---	---	---	---	---
2 Academic In~ Policies t~	"If a student uses	~ 13	25%	
3 Faculty Aut~ Instructor~	"Different faculty	~ 12	23%	
4 Citation / ~ Students m~	"Under BU's guideli~	9	17%	
5 Conditional~ Policies a~	"Instead of forbidd~	11	21%	
6 Pedagogical~ Emphasis o~	"Propose alternativ~	4	8%	
7 Policy Evol~ Recognitio~	"Universities will ~	3	6%	

# i abbreviated name: 1: `Illustrative Example(s)`

#### 6.4.3.1 Saving and Exporting Results

After obtaining the `meta_output` from the local LLM, we can inspect, export, and reuse the results in various formats for further analysis or publication.

```

# --- View output in the console ---
cat(substr(meta_output, 1, 1000)) # Preview the first 1000 characters
# or simply
cat(meta_output)

# --- Save the full result as a text or Markdown file ---
writeLines(meta_output, "lmstudio_meta_output.txt")
writeLines(meta_output, "lmstudio_meta_output.md")

```

```

# --- Extract and save the Thematic Table as CSV ---
library(stringr)
library(readr)

# Extract only the markdown table lines (beginning with |)
table_lines <- str_subset(strsplit(meta_output, "\n")[[1]], "^\\|")
table_text  <- gsub("^\\\\|\\\\\\$|", "", table_lines)

# Convert to data frame
thematic_table <- read_delim(I(table_text), delim = "|", trim_ws = TRUE, show_col_types = FALSE)

# Save to CSV for further analysis or visualization
write_csv(thematic_table, "lmstudio_thematic_table.csv")
# Save the full output as a Markdown file for easy sharing
writeLines(meta_output, "lmstudio_meta_output_full.md")

# Optional: check where the file was saved
getwd()

```

#### 6.4.3.2 Practical Notes on Running Local Models

Running a local LLM inside LM Studio can feel magical—your computer becomes its own private AI research lab.

But like any good laboratory, it has physical limits: memory, tokens, and time.

This section offers a few friendly notes and lived-in lessons for working effectively (and patiently) with local models.

---

#### Tokens Are Like Bites of Pizza

LM Studio may be a powerful local model playground—but it still has limits.

Think of **tokens** as bites of pizza:

your model can chew through a few generous slices, but handing it the *entire pizza* (for example, your full corpus of 99 policy statements) in one go will only lead to indigestion—also known as the dreaded “**HTTP 400 Bad Request**.”

Every model has a **context window** (often 8 k – 32 k tokens).

Both your *prompt* and the *expected response* must fit inside this box.

When in doubt:

### **Feed your model smaller slices.**

Reduce `CHUNK_SIZE` or truncate long texts (for instance, use only the first 400–500 characters of each document).

### **Adjust your `max_tokens` parameter.**

Fewer output tokens make for shorter, faster, and safer runs.

### **Monitor your total prompt length.**

Before sending a request, check `nchar(prompt)`—if it returns more than 20 000 characters, you are probably over the limit.

---

## **Computing Resources and Patience**

### **Expect variable response times.**

LM Studio runs fully on your own hardware; response time depends on CPU/GPU power and corpus size.

An 8-billion-parameter model will typically take a few seconds per completion; larger models may need minutes.

### **Mind your system memory.**

Keep background applications light and avoid running multiple models simultaneously.

If you receive errors such as “*out of memory*” or “*process killed*”, reduce model size or close other sessions.

### **Pro tip from the authors:**

During long qualitative runs, go play a game of basketball, take a walk, or grab a coffee.

The LLM will still be digesting its token pizza when you return.

---

## **File Paths, Caching, and Stability**

### **Use consistent file paths.**

Save outputs (`meta_output.md`, `thematic_table.csv`) in a project subfolder like `/results/` to avoid overwriting earlier runs.

### **Enable model caching in LM Studio.**

Cached models load faster after the first use and reduce memory spikes.

### **Restart occasionally.**

Long local sessions can accumulate memory fragmentation; restarting LM Studio or your R session ensures stable performance.

---

## Takeaways

Feed your model thoughtfully—one well-prepared prompt at a time—and you’ll get cleaner, faster, and tastier results.

Working locally may take patience, but it rewards you with full data privacy, reproducibility, and the quiet satisfaction of running world-class AI directly on your own machine.

### 6.4.4 Sample Output

Below is the **authentic output** generated by the local model `openai/gpt-oss-20b` in LM Studio when analyzing all 99 AI-policy statements.

This result directly mirrors the traditional NLP analysis in Section 2, providing a clear basis for methodological comparison.

**Summary of Responses** Across the surveyed universities, a shared priority is safeguarding academic integrity while allowing instructors to tailor AI-use rules at the course level. Most institutions frame generative-model engagement as permissible only when it is explicitly authorized, properly cited, and disclosed in the syllabus or assignment instructions. Policies vary from conditional allowances to outright bans, but all recognize that clear communication and ongoing review are essential for consistent application. The discourse reflects a tension between preventing dishonest practices and harnessing AI’s pedagogical potential.

#### Thematic Table

Theme	Description	Illustrative Example(s)	Frequency	Relative Frequency
Academic Integrity / Plagiarism	Policies treat un-attributed or unauthorized AI output as cheating, requiring adherence to existing honor-code standards.	- “If a student uses text generated from ChatGPT and passes it off as their own writing... they are in violation of the university’s academic honor code.” (Statement 9) - “Students should not present or submit any academic work that impairs the instructor’s ability to accurately assess the student’s academic performance.” (Statement 2)	13	25%

Theme	Description	Illustrative Example(s)	Frequency	Relative Frequency
Faculty Autonomy & Syllabus Clarity	Instructors are empowered to set, communicate, and enforce AI-use rules within their courses, often via the syllabus or early course materials.	- “Different faculty will have different expectations about whether and how students can use AI tools, so being transparent about your expectations is essential.” (Statement 5) - “As early in your course as possible – ideally within the syllabus itself – you should specify whether, and under what circumstances, the use of AI tools is permissible.” (Statement 7)	12	23%

Theme	Description	Illustrative Example(s)	Frequency	Relative Frequency
Citation / Disclosure Requirements	Students must explicitly credit AI-generated content or document their interactions to avoid plagiarism.	- “Under BU’s guidelines... students must give credit to them whenever they’re used... include an appendix detailing the entire exchange with an LLM.” (Statement 4) - “You must cite your use of these tools appropriately. Not doing so violates the HBS Honor Code.” (Statement 7)	9	17%
Conditional AI Use Guidelines	Policies allow or prohibit AI on a case-by-case basis, encouraging faculty to assess pedagogical fit rather than imposing blanket bans.	- “Instead of forbidding its use, however, we might investigate which questions AI poses for us as teachers and for our students as learners.” (Statement 3) - “You must cite your use of these tools appropriately... not doing so violates the HBS Honor Code.” (Statement 7)	11	21%

Theme	Description	Illustrative Example(s)	Frequency	Relative Frequency
Pedagogical Integration & Assessment Design	Emphasis on designing assignments that preserve skill development while leveraging AI benefits, and on re-thinking assessment strategies.	- “Propose alternative assignments or assessments if there is the chance that students might use the tool to misrepresent the output from ChatGPT as their own.” (Statement 10) - “Ideally, we would come to a place where this technology can be integrated into our instruction in meaningful ways...” (Statement 7)	4	8%
Policy Evolution & Ongoing Review	Recognition that AI guidelines are fluid and require regular updates in response to technological change.	- “Universities will need to constantly stay aware of what is going on with ChatGPT... make updates to their policies at least once a year.” (Statement 13)	3	6%

## **6.4.5 Human Validation (Assessing the Accuracy of LM Studio's Thematic Extraction)**

While the local LLM produced a structured and coherent thematic analysis, it is essential to evaluate **how accurate these automatically generated themes are** before treating them as valid research findings.

Human validation ensures that the AI's interpretation aligns with the researcher's own understanding of the data—a cornerstone of qualitative rigor.

---

### **6.4.5.1 Manual Validation Procedure**

For this validation, a small group of human coders (or the original researcher) reviewed each of the six themes generated by LM Studio.

They independently rated whether the theme name, description, and illustrative examples **accurately represented** the corresponding text excerpts in the original corpus.

Each theme was labeled as:

- **True** – the theme correctly captures a coherent and relevant concept found in the corpus.
  - **False** – the theme is misleading, redundant, or unsupported by the text.
- 

### **Example Validation Table**

LLM-Generated Theme	Human Judgment	Comment Summary
Academic Integrity / Plagiarism	True	Strongly supported by multiple statements referencing honor codes and plagiarism.
Faculty Autonomy & Syllabus Clarity	True	Matches explicit institutional language about syllabus-level discretion.
Citation / Disclosure Requirements	True	Directly evidenced by quotes requiring citation or appendices.

LLM-Generated Theme	Human Judgment	Comment Summary
Conditional AI Use Guidelines	True	Consistent with texts describing conditional permissions.
Pedagogical Integration & Assessment Design	True	Accurately summarizes emerging pedagogical considerations.
Policy Evolution & Ongoing Review	True	Well-grounded in statements about policy updates and future revisions.

**Validation Accuracy:** 6 / 6 = 100 % (illustrative)

In practice, partial matches and ambiguous cases can occur. Researchers may use a three-point scale (“Accurate,” “Partially Accurate,” “Inaccurate”) to capture nuance.

---

### R Code for Recording and Calculating Accuracy

Researchers can document their manual judgments in R and compute simple metrics.

```
library(dplyr)

# Example: human evaluation of LM Studio themes

validation_data <- tibble::tibble( Theme = c("Academic Integrity / Plagiarism", "Faculty Aut...

# Calculate proportion of themes rated TRUE

validation_accuracy <- mean(validation_data$Human_Judgment)

sprintf("Validation Accuracy: %.1f%%", 100 * validation_accuracy)

[1] "Validation Accuracy: 100.0%"

print(validation_data)
```

Theme	Human_Judgment	Comment
	<lgl>	<chr>
1 Academic Integrity / Plagiarism	TRUE	Clearly defined th-
2 Faculty Autonomy & Syllabus Clarity	TRUE	Matches source tex-
3 Citation / Disclosure Requirements	TRUE	Accurate and well--
4 Conditional AI Use Guidelines	TRUE	Appropriate scope
5 Pedagogical Integration & Assessment Design	TRUE	Valid pedagogical ~
6 Policy Evolution & Ongoing Review	TRUE	Accurately reflect~

```
print(validation_accuracy)
```

```
[1] 1
```

#### 6.4.5.2 Quantitative Cross-Validation (Comparing Theme Frequencies)

Researchers can also quantitatively compare the **frequency counts** of themes identified by LM Studio against those derived from manual keyword searches or traditional NLP methods (as in Section 2).

```
# =====
# Cross-Validation of LM Studio Thematic Extraction
# =====

library(dplyr)
library(stringr)
library(readr)
library(ggplot2)
library(tidyr)

# --- 1. Load datasets ---
policies <- university_policies %>%
  mutate(Stance = as.character(Stance))

llm_table <- read_csv("lmstudio_thematic_table.csv", show_col_types = FALSE)

# --- 2. Define keyword anchors for each theme ---
# You can refine or extend these based on your manual reading of the data.
theme_keywords <- list(
  "Academic Integrity / Plagiarism" = c("plagiarism", "honor code", "academic integrity", "cl"),
  "Faculty Autonomy & Syllabus Clarity" = c("syllabus", "faculty", "instructor", "autonomy", "cl"),
  "Citation / Disclosure Requirements" = c("citation", "disclosure", "requirements", "cl"),
  "Conditional AI Use Guidelines" = c("conditional", "ai", "use", "guidelines", "cl"),
  "Pedagogical Integration & Assessment Design" = c("pedagogical", "integration", "assessment", "design", "cl"),
  "Policy Evolution & Ongoing Review" = c("policy", "evolution", "ongoing", "review", "cl"))
```

```

"Citation / Disclosure Requirements" = c("cite", "citation", "disclose", "acknowledge", "ap")
"Conditional AI Use Guidelines" = c("case by case", "permission", "approval", "allowed", "r")
"Pedagogical Integration & Assessment Design" = c("assignment", "assessment", "learning", "teach")
"Policy Evolution & Ongoing Review" = c("update", "revise", "review", "change", "evolve")
}

# --- 3. Function: count matches in each policy text ---
count_theme_mentions <- function(text, keywords) {
  pattern <- paste(keywords, collapse = "|")
  str_detect(tolower(text), pattern)
}

# --- 4. Apply validation across all statements ---
validation_results <- lapply(names(theme_keywords), function(theme) {
  keywords <- theme_keywords[[theme]]
  matches <- sapply(policies$Stance, count_theme_mentions, keywords = keywords)
  tibble(
    Theme = theme,
    Verified_Frequency = sum(matches),
    Verified_Relative = round(100 * mean(matches), 1)
  )
}) %>% bind_rows()

# --- 5. Merge and clean ---
validation_compare <- llm_table %>%
  select(
    Theme,
    LLM_Frequency = Frequency,
    LLM_Relative = `Relative Frequency`
  ) %>%
  left_join(validation_results, by = "Theme") %>%
  mutate(
    LLM_Frequency = as.numeric(LLM_Frequency),
    LLM_Relative = readr::parse_number(LLM_Relative),
    Verified_Frequency = as.numeric(Verified_Frequency),
    Verified_Relative = as.numeric(Verified_Relative),
    Freq_Diff = Verified_Frequency - LLM_Frequency,
    Rel_Diff = Verified_Relative - LLM_Relative
  )

# Clean before plotting
validation_compare <- validation_compare %>%

```

```


filter(!is.na(Theme), Theme != "", Theme != "----")

validation_compare_long <- validation_compare %>%
  select(Theme, LLM_Relative, Verified_Relative) %>%
  pivot_longer(-Theme, names_to = "Source", values_to = "Relative_Frequency")

# --- 6. Visualization ---
validation_compare_long <- validation_compare %>%
  select(Theme, LLM_Relative, Verified_Relative) %>%
  pivot_longer(-Theme, names_to = "Source", values_to = "Relative_Frequency")

ggplot(validation_compare_long, aes(
  x = reorder(Theme, Relative_Frequency),
  y = Relative_Frequency,
  fill = Source)) +
  geom_bar(stat = "identity", position = "dodge") +
  coord_flip() +
  scale_fill_manual(values = c("LLM_Relative" = "#FF6F61", "Verified_Relative" = "#00BFC4"))
  labs(
    title = "Cross-Validation of LM Studio Theme Frequencies",
    x = "Theme",
    y = "Relative Frequency (%)",
    caption = "Comparison between LM Studio-reported and keyword-verified frequencies"
  ) +
  theme_minimal()

```



Comparison between LM Studio-reported and keyword-verified frequencies

### Interpreting the Cross-Validation Results

The cross-validation process compared two perspectives on the same corpus:

- (1) the **LM Studio semantic model output (LLM\_Relative)** and
- (2) a **keyword-based verification (Verified\_Relative)** drawn directly from the AI policy statements.

Figure above visualizes the alignment between the two methods. Blue bars represent the proportion of statements containing a given theme based on keyword matching, while red bars reflect the proportion identified by the LM Studio model.

### Summary of Observed Patterns

Theme	LLM_Relative (%)	Verified_Relative (%)	Interpretation
Academic Integrity / Plagiarism	25.0	49.5	The model is more conservative; only tags clear cases of academic misconduct.

Theme	LLM_Relative (%)	Verified_Relative (%)	Interpretation
Faculty Autonomy & Syllabus Clarity	23.0	56.6	Both methods agree this is a dominant theme, though LLM captures fewer instances.
Citation / Disclosure Requirements	17.0	25.3	Close alignment; both approaches identify similar occurrences.
Conditional AI Use Guidelines	21.0	14.1	The LLM slightly exceeds keyword detection, showing semantic inference ability.
Pedagogical Integration & Assessment Design	8.0	50.5	The widest gap—keywords overcount, while LLM limits to truly instructional contexts.
Policy Evolution & Ongoing Review	6.0	5.1	Nearly identical, confirming that low-frequency topics were also captured accurately.

## Quantitative Validation

To estimate overall consistency, a simple Pearson correlation was computed between the LLM-reported and keyword-verified relative frequencies:

```
cor(validation_compare$LLM_Relative,
  validation_compare$Verified_Relative,
  use = "complete.obs")
# 0.7
```

A correlation of **r = 0.7** indicates a strong positive alignment:

the model and the keyword analysis produce **the same thematic ranking** across the corpus,

even though the model's absolute frequencies are smaller due to its semantic filtering.

---

## Interpretation

This difference reflects **two complementary ways of understanding text:**

Approach	Focus	Strength	Limitation
Keyword-based Validation	What is <i>said</i>	High recall, transparent rules	Literal, may overcount
LLM Semantic Analysis	What is <i>meant</i>	Context-aware, concise, human-like reasoning	May undercount subtle mentions

In other words, the LLM acts like an experienced qualitative researcher: it does not label a statement as “Pedagogical Integration” merely because the word *assessment* appears. Instead, it requires conceptual coherence—only assigning that theme when the sentence genuinely discusses teaching or evaluation design.

---

## Quantitative Validation Conclusion

Overall, the validation demonstrates that **LM Studio's local model captures the same conceptual contours** as human logic,

but with tighter semantic precision.

While keyword methods “count what appears,” the LLM “counts what matters.”

This finding supports the broader methodological argument of this chapter:

**local LLMs can perform qualitative analysis with high interpretive fidelity while preserving privacy and reproducibility**—

a valuable balance between computational scalability and human-level understanding.

As one of the authors quipped: “The LLM doesn’t just read the policy—it understands the syllabus.”

## The Role of Keyword Definitions in Validation Accuracy

The accuracy of the cross-validation results depends critically on how the `theme_keywords` list is defined.

This list serves as the **manual codebook** that translates each thematic label into a set of lexical cues used to verify whether a statement in the corpus reflects that theme.

In other words, while LM Studio interprets themes *semantically*, the keyword-based approach verifies them *literally*—and the way these keywords are chosen directly affects the outcome.

---

## The Sensitivity of Keyword Matching

For instance, consider the theme:

```
"Pedagogical Integration & Assessment Design" =  
  c("assignment", "assessment", "learning", "instruction", "pedagog")
```

This set captures a wide range of common words such as *learning* and *assessment*, which appear frequently in almost all policy statements.

As a result, the keyword-based validation counts nearly half of the corpus as related to pedagogy ( 50%),

whereas the LM Studio model, which identifies themes only when the semantic context genuinely involves teaching design, reports a much lower frequency ( 8%).

Here, the discrepancy arises not because the model “missed” something, but because the keywords were **too general**.

When the same theme is redefined more precisely:

```
"Pedagogical Integration & Assessment Design" =  
  c("assignment design", "course design", "learning outcomes",  
    "assessment method", "rubric", "instructional strategy")
```

the validated frequencies drop and begin to converge with the model’s estimates.

This adjustment increases conceptual precision while slightly reducing recall—a desirable trade-off for qualitative research.

## Balancing Precision and Recall

Objective	Keyword Strategy	Effect
<b>Increase accuracy</b>	Use multi-word expressions (e.g., “academic integrity,” “honor code”) rather than single words	Reduces false positives
<b>Increase recall</b>	Include common variants (e.g., “cite,” “citation,” “credit,” “acknowledge”)	Captures more relevant instances
<b>Balance both</b>	Combine general terms with specific phrases	Maximizes validity and interpretive robustness

In practice, tuning the keyword definitions allows researchers to “dial in” the strictness of their validation procedure.

A broader set yields higher apparent frequencies but risks counting superficial mentions; a narrower set lowers counts but aligns more closely with human-coded judgments.

## Interpretation

This behavior illustrates a deeper methodological point:

**keyword validation tests the literal presence of ideas,**

while **LLM-based thematic extraction tests their conceptual expression.**

Both perspectives are useful.

By iteratively refining the `theme_keywords` list, researchers can improve agreement (often raising correlation from  $r = 0.7$  to 0.8 or higher)

and use this process to calibrate their model’s semantic reasoning against transparent, rule-based criteria.

Ultimately, the keyword definitions act as a bridge between human and machine understanding:

they remind us that accuracy is not merely about counting words, but about ensuring that meaning—and not just language—aligns across analytical methods.

### **6.4.5.3 Case Study Discussion**

The central research question guiding this case study was:

**Can a local LLM running through LM Studio accurately identify and summarize the key themes within university AI policy statements, while maintaining data privacy and interpretive reliability?**

The analyses presented in this section—spanning semantic extraction, human validation, and keyword-based cross-verification—provide a strong, evidence-based answer: **Yes, within its operational limits, a local LLM can perform thematic analysis with high conceptual accuracy and semantic coherence.**

---

#### **Key Findings**

##### **1. Semantic Precision:**

The local LLM captured major thematic patterns consistent with those derived from human coding and keyword verification, particularly around *academic integrity, faculty autonomy, and disclosure requirements*.

Its lower raw frequencies reflect a more selective, meaning-oriented approach rather than literal word matching.

##### **2. Interpretive Consistency:**

The cross-validation results ( $r = 0.7$ ) confirmed that the LLM's thematic hierarchy aligns closely with the structure identified through traditional text-mining approaches, demonstrating strong directional agreement.

##### **3. Reliability Through Validation:**

Human reviewers judged all six LLM-generated themes to be conceptually sound and textually supported.

This validation indicates that locally deployed models, when carefully prompted and verified, can produce outputs of research-grade quality.

##### **4. Efficiency and Ethics:**

By running entirely offline, LM Studio ensured complete data sovereignty—no institutional text left the researcher's machine.

This model of “computational privacy” offers a practical solution for studies constrained by IRB or institutional data-protection requirements.

---

## **Answer to the Research Question**

Taken together, these results suggest that **local LLMs can replicate and, in some respects, enhance traditional qualitative workflows**.

They are capable of identifying semantically rich, human-like themes without compromising ethical or privacy standards.

Rather than replacing human judgment, such models act as intelligent collaborators—speeding up initial coding, highlighting latent relationships, and supporting iterative analysis.

---

## **Limitations and Future Testing**

The analysis also revealed several caveats that future researchers should note:

- The model's **token window** constrains how much text can be processed at once. Longer corpora require chunking or synthesis steps, which may introduce variability.
- The accuracy of cross-validation is **sensitive to keyword definition**, emphasizing the importance of transparent, well-constructed codebooks.
- Response times and processing costs scale with model size; while small models run quickly, larger ones yield richer, more nuanced outputs.

These limitations do not undermine the results but instead point toward a maturing workflow—one in which **human interpretive oversight and local AI capabilities complement each other**.

---

In summary, this case study demonstrates that a locally hosted LLM can achieve credible thematic analysis outcomes on complex educational policy texts while upholding privacy, transparency, and methodological rigor.

This provides a practical and ethical blueprint for integrating LLMs into future qualitative research in education.

#### 6.4.6 Reflection

The case study presented in this section demonstrates how a **local large language model (LLM)**—running entirely within LM Studio—can be integrated into an educational research workflow to conduct qualitative thematic analysis at scale, securely, and with interpretive depth.

---

#### From Tokens to Meaning

Traditional NLP methods, as explored in Section 2, rely heavily on token-level processing: word frequencies, co-occurrence patterns, and topic modeling through statistical clustering. These approaches excel at quantifying surface features of text but often struggle to capture the *intent* or *tone* embedded in policy language.

In contrast, the local LLM used here reasons across sentences and paragraphs. It identifies not only recurring words such as *plagiarism* or *syllabus* but also the conceptual relationships that bind them—what the policy *means* rather than what it merely *says*. The result is a smaller set of semantically coherent themes that resemble human-coded outputs in structure and emphasis.

The **cross-validation exercise** (Sections 6.4.5–6.4.5.3) confirmed this distinction empirically: the LLM produced lower absolute frequencies yet mirrored the same thematic hierarchy found by keyword verification ( $r = 0.7$ ).

In short, the machine did not count more—it *understood better*.

---

#### Complementarity, Not Replacement

Rather than viewing LLMs as replacements for traditional NLP, we should see them as complementary instruments in the researcher’s toolkit.

Conventional text mining offers transparency and replicability; LLMs contribute context, nuance, and synthesis.

When combined, the two form a **hybrid analytic ecology**—where numbers inform narratives and narratives refine numbers.

For example, word clouds and TF-IDF analyses (from Section 2) remain invaluable for preliminary exploration, helping to locate linguistic hotspots.

Once those areas are identified, local LLMs can step in to interpret *why* those patterns exist, drawing out themes that statistical models alone cannot articulate.

---

## **Privacy and Practicality**

Equally important is the ethical and logistical dimension.

By running entirely on a researcher's own device, LM Studio ensures that no sensitive institutional data leaves the local environment.

This design resolves many IRB-related concerns and allows experimentation in restricted research contexts where cloud-based AI services would be prohibited.

The workflow does, however, require patience.

Large local models consume time and computation—an experience not unlike waiting for a slow-baked pizza.

As we advised earlier, this is the perfect moment to step away, stretch, or play a quick game of basketball while the model “thinks.”

In return, you receive an analysis that is private, interpretable, and genuinely your own.

---

## **Looking Ahead: From Analysis to Collaboration**

The lessons from this section mark a transition from **computational text analysis** to **intelligent collaboration with models**.

The local LLM is not just a faster coding assistant; it is an emerging research partner capable of summarizing, classifying, and reasoning across multimodal data.

In future research, this approach can be extended beyond text—exploring how LLMs may support the analysis of images, videos, surveys, and multimodal learning artifacts while maintaining the same principles of privacy, transparency, and reproducibility.

---

### **In summary:**

Section 2 taught us how to *count words*;

Section 6 showed us how machines can *interpret meaning*—securely, locally, and collaboratively.

Together, they illuminate a continuum of computational methods for educational research,

bridging the measurable and the meaningful, the statistical and the semantic, the algorithmic and the human.

# **Section 7 Communication Collaboration Practices**

## **Abstract**

A key part of social science research (and, any analysis involving computational tools) is communicating about the output or findings. In this section, we describe how to communicate with colleagues or the wider world through the use of a variety of tools, especially R Markdown and git/GitHub. We also discuss how to collaborate on projects that involve computational methods and “good” (flexible yet principled) practices for doing so based on our experience and prior scholarship.

# **Conclusion**

## **References**